

The Sources of Capital Misallocation[†]

By JOEL M. DAVID AND VENKY VENKATESWARAN*

We develop a methodology to disentangle sources of capital “misallocation,” i.e., dispersion in value-added/capital. It measures the contributions of technological/informational frictions and a rich class of firm-specific factors. An application to Chinese manufacturing firms reveals that adjustment costs and uncertainty, while significant, explain only a modest fraction of the dispersion, which stems largely from other factors: a component correlated with productivity and a fixed effect. Adjustment costs are more salient for large US firms, though other factors still account for the bulk of the dispersion. Technological/markup heterogeneity explains a limited fraction in China, but a potentially large share in the United States. (JEL D22, D24, D25, E22, G31, L60, O11, O14, O47, P31)

A large and growing body of work analyzes the “misallocation” of productive resources across firms, usually measured by dispersion in the static average products of inputs (e.g., value-added/input ratios), and the resulting adverse effects on aggregate productivity and output. A number of recent studies examine the role of specific factors hindering period-by-period equalization of input productivity ratios. Examples of such factors include adjustment costs, imperfect information, financial frictions, as well as firm-specific “distortions” stemming from economic policies or other institutional features. The importance of disentangling the role of these forces is self-evident. For one, a central question, particularly from a policy standpoint, is whether observed variation in input products stems largely from efficient sources, e.g., technological factors like adjustment costs or heterogeneity in production technologies, or inefficient ones, such as policy-induced distortions or markups. Similarly, understanding the exact nature of distortions, e.g., the extent to which they are correlated with firm characteristics, is essential to analyze their

*David: Department of Economics, University of Southern California, Los Angeles, CA 90089 (email: joeldavi@usc.edu); Venkateswaran: Federal Reserve Bank of Minneapolis and NYU Stern School of Business, Department of Economics, New York, NY 10012 (email: vvenkate@stern.nyu.edu). Mikhail Golosov was the coeditor for this article. A previous version of this paper was circulated under the title “Capital Misallocation: Frictions or Distortions?” We thank the editor and anonymous referees, Pete Klenow, Diego Restuccia, Francisco Buera, Daniel Xu, Richard Rogerson, Matthias Kehrig, Loukas Karabarbounis, Virgiliu Midrigan, Andy Atkeson, Hugo Hopenhayn, and Russ Cooper for their helpful suggestions, Nicolas Petrosky-Nadeau and Oleg Itshkoki for insightful discussions, Junjie Xia for assistance with the Chinese data, and many seminar and conference participants for useful comments. The authors declare that they have no relevant or material financial interests that relate to the research described in this paper.

[†]Go to <https://doi.org/10.1257/aer.20180336> to visit the article page for additional materials and author disclosure statements.

implications beyond static misallocation, for example, on firm entry and exit decisions and investments that influence future productivity.¹

In this paper, we develop and implement a tractable methodology to distinguish various sources of dispersion in average revenue products of capital (*arpk*) using observable data on value-added and inputs. Our analysis proceeds in two steps. First, we augment a standard general equilibrium model of firm dynamics with a number of forces that contribute to ex post dispersion in the static *arpk*, specifically (i) capital adjustment costs; (ii) informational frictions, in the form of imperfect knowledge about firm-level fundamentals (e.g., productivity or demand); and (iii) other firm-specific factors, meant to capture all other forces influencing investment decisions, including unobserved heterogeneity in markups and/or production technologies, financial frictions, or institutional/policy-related distortions. In this first part of our analysis, rather than take a stand on the exact nature of these factors, we adopt a flexible specification that allows for time-variation and correlation with firm characteristics. The environment is an extension of the canonical Hsieh and Klenow (2009) framework to include dynamic considerations in firms' investment decisions. The main contribution of this part is an empirical strategy that precisely measures the contribution of each force to observed *arpk* dispersion using widely available firm-level data.

In the second part of our analysis, we explore various candidates for the firm-specific factors in (iii). First, we extend our methodology to investigate the extent to which the observed dispersion in *arpk* could stem from unobserved heterogeneity in markups and production technologies. Next, we analyze policies that restrict the size of firms and study a model of financial/liquidity considerations. We show how these two forces can manifest themselves as firm-specific factors similar to those considered in the first part.

Our key innovation is to explore the sources of *arpk* dispersion within a unified framework and thus provide a more robust decomposition. In contrast, focusing on particular sources while abstracting from others, a common approach in the literature, is potentially problematic. When the data reflect the combined influence of a number of factors, examining them one-by-one can lead to biased assessments of their severity and contributions to the observed dispersion.

To understand the measurement difficulty, consider, as an example, convex adjustment costs. When they are the only force present, there is an intuitive, one-to-one mapping to a single moment, e.g., investment variability: the more severe the adjustment friction, the less volatile is investment. Now, suppose that there are other factors that also dampen investment volatility (e.g., a distortion correlated with productivity or size). In this case, using the variance of investment alone to draw inferences about adjustment costs leads to an upward bias. As a second example, consider the effects of firm-level uncertainty, which reduces the contemporaneous correlation between investment and productivity. However, a low correlation could also be the result of other firm-specific factors (e.g., markups) that are uncorrelated with productivity, making this single moment an inadequate measure of the quality of information.

¹ See Restuccia and Rogerson (2017) for an in-depth discussion of these margins.

Our strategy for disentangling these forces is based on a simple insight: although each moment is a complicated function of multiple factors, making any single one insufficient for identification, combining the information in a wider set of moments can be extremely helpful in disentangling these factors. Indeed, we show that allowing these forces to act in tandem is essential to reconcile a broad set of moments from the covariance matrix of firm-level investment and value-added. We formalize this intuition using a tractable special case, when firm-level productivity follows a random walk. In this case, we derive analytic expressions for the moments and prove that a set of four carefully chosen moments, namely, (i) the variance of investment, (ii) the autocorrelation of investment, (iii) the correlation of investment with past productivity, and (iv) the covariance of *arpk* with productivity together uniquely identify adjustment costs, uncertainty and the magnitude and correlation structure of other firm-specific factors.

The intuition behind this result is easiest to see in a simple pairwise analysis. As an example, consider the challenge described earlier of disentangling adjustment costs from other idiosyncratic factors that dampen the response of investment to productivity. Both forces depress the variability of investment. However, they have opposing effects on its autocorrelation: convex adjustment costs create incentives to smooth investment over time and so increase its serial correlation. A distortion that reduces the responsiveness to productivity, on the other hand, raises the relative weight of other, more transitory considerations in the investment decision, lowering the serial correlation. Thus, holding all else fixed, these two moments allow us to separate the two forces. Similar arguments can be developed for the remaining factors as well. In our quantitative work, where we depart from the polar random walk case, we demonstrate numerically that the same intuition carries through.

This logic also underlies the second part of our analysis, where we dig deeper into factors other than adjustment/information frictions. For example, we use moments of labor and materials usage to investigate the role of unobserved heterogeneity in markups and technologies (specifically, capital elasticities). Under the assumption that the choice of materials is distorted only by market power, markup dispersion is pinned down by the dispersion in materials' share of revenues. Technology dispersion can be bounded from above using the observed covariance between the average products of capital and labor. Intuitively, holding returns to scale fixed, a high production elasticity of capital implies a low labor elasticity, so this type of heterogeneity is a source of negative covariance between capital and labor products. The more positively correlated these are in the data, the lower the scope for *arpk* dispersion from this channel.

We apply our methodology to data on manufacturing firms in China over the period 1998–2009. We find that adjustment and informational frictions play economically significant roles in influencing observed investment dynamics. However, they account for only a relatively modest fraction of *arpk* dispersion among Chinese firms, about 1 and 10 percent, respectively, leading to losses in aggregate total factor productivity (TFP) of 1 and 8 percent (relative to the undistorted first-best). This implies that a substantial portion of *arpk* dispersion in China is due to other firm-specific factors. In particular, we find a large role for factors correlated with productivity and ones that are essentially permanent. These account for about 47 and 44 percent of overall *arpk* dispersion, respectively, leading to TFP losses of 38

and 36 percent.² These findings are driven in large part by two observations: first, firm-level investment is neither extremely volatile nor highly serially correlated. The latter bounds the potential for convex adjustment costs, which create incentives to smooth investment over time. In combination with the former, this leads us to ascribe a large role to correlated distortions, which reduce investment volatility without increasing the serial correlation. Importantly, as we discuss below, these insights continue to hold even when we introduce nonconvexities in the adjustment cost function. Second, uncertainty over future productivity, while significant, is simply not large enough to account for the majority of *arpk* dispersion observed in the data.

We also apply the methodology to data on publicly traded firms in the United States. Although the two sets of firms are not directly comparable, the US numbers serve as a useful benchmark to put our results for China in context.³ As one would expect, the overall degree of *arpk* dispersion is considerably smaller for publicly traded US firms. More interestingly, a larger share (about 11 percent) of the observed dispersion is accounted for by adjustment costs, which depress aggregate TFP by about 2 percent. Uncertainty plays a smaller role than among Chinese firms, as do other correlated factors: these account for about 7 and 14 percent of overall *arpk* dispersion, respectively, reducing aggregate TFP by 1 and 3 percent. However, even for these firms, firm-specific fixed factors, although considerably smaller in absolute magnitude than in China, generate a large share of the observed dispersion in *arpk*, accounting for about 65 percent of the total, with associated TFP losses of 13 percent. In other words, even in the United States, factors other than technological and informational frictions play a significant role in determining capital allocations.

What are these firm-specific factors? First, we find modest scope for unobserved variation in markups or production technologies in China: together, they account for at most 27 percent of *arpk* dispersion (4 and 23 percent, respectively). Intuitively, we do not see much variation in materials' share of revenues in China, suggesting only small markup dispersion, and the average products of labor and capital are highly correlated, limiting the potential for heterogeneity in capital intensities. In contrast, for US publicly traded firms, variation in markups/technologies can explain as much as 58 percent of *arpk* dispersion (14 and 44 percent, respectively). These results suggest that unobserved heterogeneity is a promising explanation for much of the observed "misallocation" in the United States, but that the predominant drivers among Chinese firms lie elsewhere, e.g., additional market frictions or institutional/policy-related distortions. For example, we show that our estimates of size/productivity-dependent factors could be picking up the effects of size-dependent government policies and certain forms of financial market imperfections. However, disentangling these two forces from other sources of correlated factors requires data beyond value-added and inputs (e.g., firm-level financial data).

We show that these patterns, in particular, the contributions of the various forces to observed *arpk* dispersion, are robust to a number of variations of our baseline setup.

²Our method also allows for distortions that are transitory and uncorrelated with firm characteristics. However, our estimation finds them to be negligible.

³We also report results for Chinese publicly traded firms as well as Colombian and Mexican manufacturing firms. The results regarding the role of various factors in driving observed dispersion are quite similar to our baseline findings for Chinese manufacturers.

First, they are largely unchanged when we allow for nonconvex adjustment costs. The main insight that underlies our baseline estimates emerges here as well: examining moments in isolation can paint a distorted picture of the forces driving investment dynamics. For example, the low serial correlation of investment in the data by itself might seem to indicate large nonconvexities, but high fixed costs of adjustment also tend to make firm-level investment more volatile and induce substantial “inaction” (i.e., periods with little or no investment), patterns which are inconsistent with the data. Models with only adjustment costs, even those with sophisticated specifications, struggle to simultaneously match these patterns and can produce very different estimates depending on the choice of target moments. For example, targeting investment variability and inaction (and ignoring the serial correlation, as in, e.g., Asker, Collard-Wexler, and De Loecker 2014) results in much larger estimates for convex costs relative to a strategy of targeting the serial correlation (and ignoring variability/inaction, as in, e.g., Cooper and Haltiwanger 2006), which instead produces substantial fixed costs. This underscores the value of a strategy like ours, which targets a broad set of moments with a more flexible model. Our estimation reconciles these seemingly inconsistent data patterns by ascribing an important role to other firm-specific factors, particularly when it comes to generating *arpk* dispersion.

Next, we show that distortions in the labor choice do not alter our main conclusions. A version of our model in which labor is subject to the same frictions and distortions as capital leads to a very similar decomposition of *arpk* dispersion. However, since both inputs are affected by each of the forces, the associated implications for aggregate TFP are much larger. Lastly, we show that our main results remain valid across a number of robustness exercises aimed at addressing measurement-related issues, sectoral heterogeneity, and parameter choices.

The paper is organized as follows. Section I describes our model of frictional investment. Section II spells out our identification strategy using the analytically tractable random walk case, while Section III details our numerical analysis and presents our quantitative results. Section IV further investigates the potential sources of firm-specific idiosyncratic factors. Section V explores a number of variants on our baseline approach. We summarize our findings and discuss directions for future research in Section VI.

Related Literature.—Our paper relates to several branches of literature. We bear a direct connection to the large body of work measuring and quantifying the effects of resource misallocation.⁴ Following the seminal work of Hsieh and Klenow (2009) and Restuccia and Rogerson (2008), recent attention has shifted toward analyzing the roles of specific factors. Important contributions include Asker, Collard-Wexler, and De Loecker (2014) on adjustment costs; Buera, Kaboski, and Shin (2011), Midrigan and Xu (2014), Moll (2014), and Gopinath et al. (2017) on financial frictions; David, Hopenhayn, and Venkateswaran (2016) on uncertainty; and Peters (2016) on markup dispersion. Several recent papers analyze subsets of these factors in combination. For example, Gopinath et al. (2017) study the interaction of capital adjustment costs and size-dependent financial frictions in determining the recent

⁴Restuccia and Rogerson (2017) and Hopenhayn (2014) provide recent overviews of this line of work.

dynamics of capital allocation in Spain. Kehrig and Vincent (2017) combine financial and adjustment frictions to investigate misallocation within firms, while Song and Wu (2015) estimate a model with adjustment costs, permanent distortions, and heterogeneity in markups/technologies.

Our primary contribution is to develop a unified framework that encompasses many of these factors and devise an empirical strategy based on observable firm-level data to disentangle them. Our results, both analytical and quantitative, highlight the importance of studying a broad set of forces in tandem. This breadth is partly what distinguishes us from the work of Song and Wu (2015), who abstract from time-variation in firm-level distortions (as well as in firm-specific markups/technologies), ruling out, by assumption, any role for so-called “correlated” or size-dependent distortions.⁵ Many papers in the literature (e.g., Restuccia and Rogerson 2008; Bartelsman, Haltiwanger, and Scarpetta 2013; Hsieh and Klenow 2014; and Bento and Restuccia 2017) emphasize the need to distinguish such factors from those that are orthogonal to productivity. This message is reinforced by our quantitative findings, which reveal a significant role for correlated factors (in addition to uncorrelated, permanent ones), particularly in developing countries such as China. Our modeling of these factors as implicit taxes correlated with productivity follows the approach taken by, e.g., Guner, Ventura, and Xu (2008); Restuccia and Rogerson (2008); Bartelsman, Haltiwanger, and Scarpetta (2013); Buera, Moll, and Shin (2013); Hsieh and Klenow (2014); and Buera and Fattal-Jaef (2018).

Our methodology and findings also have relevance beyond the misallocation context, notably, for studies of adjustment and informational frictions. A large literature has examined the implications of adjustment costs, examples of which include Cooper and Haltiwanger (2006), Khan and Thomas (2008), and Bloom (2009). Our analysis shows that accounting for other firm-specific factors acting on firms’ investment decisions is potentially crucial in order to accurately estimate the severity of these frictions and reconcile a broader set of micro-level moments. A similar point applies to recent work on quantifying firm-level uncertainty, for example, Bloom (2009); Bachmann and Elstner (2015); and Jurado, Ludvigson, and Ng (2015).

I. The Model

We consider a discrete time, infinite-horizon economy, populated by a representative household. The household inelastically supplies a fixed quantity of labor N and has preferences over consumption of a final good. The household discounts time at rate β . The household side of the economy is deliberately kept simple as it plays a limited role in our study. Throughout the analysis, we focus on a stationary equilibrium in which all aggregate variables remain constant.

Production.—A continuum of firms of fixed measure 1, indexed by i , produce intermediate goods using capital and labor according to

$$(1) \quad Y_{it} = K_{it}^{\hat{\alpha}_1} N_{it}^{\hat{\alpha}_2}, \quad \hat{\alpha}_1 + \hat{\alpha}_2 \leq 1.$$

⁵ We also differ from Song and Wu (2015) in our explicit modeling (and measurement) of information frictions and in our approach to quantifying heterogeneity in markups/technologies.

These intermediate goods are bundled to produce the single final good using a standard constant elasticity of substitution (CES) aggregator

$$Y_t = \left(\int \hat{A}_{it} Y_{it}^{\frac{\theta-1}{\theta}} di \right)^{\frac{\theta}{\theta-1}},$$

where $\theta \in (1, \infty)$ is the elasticity of substitution between intermediate goods and \hat{A}_{it} represents a firm-specific idiosyncratic component in production/demand. This is the only source of fundamental uncertainty in the economy (i.e., we abstract from aggregate risk).

Market Structure and Revenue.—The final good is produced frictionlessly by a representative competitive firm. This yields a standard demand function for intermediate good i ,

$$Y_{it} = P_{it}^{-\theta} \hat{A}_{it}^{\theta} Y_t \Rightarrow P_{it} = \left(\frac{Y_{it}}{Y_t} \right)^{-\frac{1}{\theta}} \hat{A}_{it},$$

where P_{it} denotes the relative price of good i in terms of the final good, which serves as numéraire. Revenues (here, equal to value-added) for firm i at time t are

$$P_{it} Y_{it} = Y_t^{\frac{1}{\theta}} \hat{A}_{it} K_{it}^{\alpha_1} N_{it}^{\alpha_2},$$

where

$$\alpha_j = \left(1 - \frac{1}{\theta} \right) \hat{\alpha}_j, \quad j = 1, 2.$$

This framework accommodates two alternative interpretations of the idiosyncratic component, \hat{A}_{it} : as a firm-specific shifter of either quality/demand or productive efficiency.

Input Choices.—In our baseline analysis, we assume that firms hire labor period-by-period under full information and in an undistorted fashion at a competitive wage, W_t .⁶ At the end of each period, firms choose capital for the following period. Investment is subject to quadratic adjustment costs, given by

$$(2) \quad \Phi(K_{it+1}, K_{it}) = \frac{\hat{\xi}}{2} \left(\frac{K_{it+1}}{K_{it}} - (1 - \delta) \right)^2 K_{it},$$

where $\hat{\xi}$ parameterizes the severity of the adjustment cost and δ is the rate of depreciation.⁷

Investment decisions are likely to be affected by a number of additional factors (other than productivity/demand and the level of installed capital). These could

⁶We relax this assumption later in the paper in two separate exercises. First, in online Appendix E.1, we introduce labor market distortions in the form of firm-specific “taxes” and show that this formulation changes the interpretation of our measure of productivity, but does not affect our identification strategy or conclusions about the sources of *arpk* dispersion. Second, in Section VB, we subject the labor choice to all the frictions that distort investment: adjustment costs, informational frictions and other distortionary factors. This setup also leads to a very similar problem with suitably redefined productivity and curvature parameters.

⁷We generalize this specification to include nonconvex costs in Section VA and show that our main quantitative results continue to hold.

originate from distortionary government policies (e.g., taxes, size restrictions or regulations, or other features of the institutional environment), from other market frictions that are not explicitly modeled (e.g., financial frictions) or from unmodeled heterogeneity in markups/production technologies. For now, we do not take a stand on the precise nature of these factors and, following, e.g., Hsieh and Klenow (2009), model them as firm-specific proportional “taxes” on the flow cost of capital. We denote these by T_{it+1}^K and, in a slight abuse of terminology, refer to them as “distortions” or wedges throughout the paper, even though they may partly reflect efficient factors (for example, heterogeneity in production functions).⁸ In Section IV, we demonstrate how progress can be made in further disentangling some of these sources.

The firm’s problem in a stationary equilibrium can be represented in recursive form as (we suppress the time subscript on all aggregate variables)

$$\begin{aligned} \mathcal{V}(K_{it}, \mathcal{I}_{it}) \\ = \max_{N_{it}, K_{it+1}} E_{it} \left[Y_{\theta}^1 \hat{A}_{it} K_{it}^{\alpha_1} N_{it}^{\alpha_2} - WN_{it} - T_{it+1}^K K_{it+1} (1 - \beta(1 - \delta)) - \Phi(K_{it+1}, K_{it}) \right] \\ + \beta E_{it} [\mathcal{V}(K_{it+1}, \mathcal{I}_{it+1})], \end{aligned}$$

where $E_{it}[\cdot]$ denotes expectations conditional on \mathcal{I}_{it} , the firm’s information set at the time it chooses K_{it+1} (described in more detail below). Note that the wedge T_{it+1}^K (which applies to $1 - \beta(1 - \delta)$, the per-unit user cost of capital) distorts both the capital decision and the capital-labor ratio. In other words, it is both a “scale” and “mix” distortion.⁹

After maximizing over N_{it} , this becomes

$$\begin{aligned} (3) \quad \mathcal{V}(K_{it}, \mathcal{I}_{it}) = \max_{K_{it+1}} E_{it} \left[GA_{it} K_{it}^{\alpha} - T_{it+1}^K K_{it+1} (1 - \beta(1 - \delta)) - \Phi(K_{it+1}, K_{it}) \right] \\ + E_{it} \beta [\mathcal{V}(K_{it+1}, \mathcal{I}_{it+1})], \end{aligned}$$

where $\alpha \equiv \frac{\alpha_1}{1 - \alpha_2}$ is the curvature of operating profits (value-added less labor expenses) and $A_{it} \equiv \hat{A}_{it}^{\frac{1}{1 - \alpha_2}}$ is the firm-specific profitability of capital. In a slight abuse of terminology, we refer to A_{it} simply as firm-specific productivity. The term $G \equiv (1 - \alpha_2) \left(\frac{\alpha_2}{W} \right)^{\frac{\alpha_2}{1 - \alpha_2}} Y_{\theta}^{\frac{1}{1 - \alpha_2}}$ is a constant that captures the effects of aggregate variables.

Equilibrium.—We can now define a *stationary equilibrium* in this economy as (i) a set of value and policy functions for the firm, $\mathcal{V}(K_{it}, \mathcal{I}_{it})$, $N_{it}(K_{it}, \mathcal{I}_{it})$, and $K_{it+1}(K_{it}, \mathcal{I}_{it})$, (ii) a wage W , and (iii) a joint distribution over $(K_{it}, \mathcal{I}_{it})$ such that (a) taking as given wages and the law of motion for \mathcal{I}_{it} , the value and policy functions solve the firm’s optimization problem, (b) the labor market clears, and (c) the joint distribution remains constant through time.

⁸The timing convention implies that the wedge T_{it+1}^K affects the firm’s choice of K_{it+1} .

⁹In Section VB, when the firm’s labor choice is assumed to be subject to the same frictions/distortions as its investment decision, the wedge is a pure scale distortion, i.e., it does not distort the capital-labor ratio.

Characterization.—We use perturbation methods to solve the model.¹⁰ In particular, we log-linearize the firm's optimality conditions and laws of motion around $A_{it} = \bar{A}$ (the unconditional average level of productivity) and $T_{it}^K = 1$ (i.e., no distortions). Online Appendix A.1 derives the following log-linearized Euler equation:¹¹

$$(4) \quad k_{it+1}((1 + \beta)\xi + 1 - \alpha) = E_{it}[a_{it+1} + \tau_{it+1}] + \beta\xi E_{it}[k_{it+2}] + \xi k_{it},$$

where ξ and τ_{it+1} are rescaled versions of the adjustment cost parameter, $\hat{\xi}$, and the distortion, $\log T_{it+1}^K$, respectively.

Stochastic Processes.—We assume that the productivity, A_{it} , follows an AR(1) process in logs with normally distributed i.i.d. innovations, i.e.,

$$(5) \quad a_{it} = \rho a_{it-1} + \mu_{it}, \quad \mu_{it} \sim \mathcal{N}(0, \sigma_\mu^2),$$

where ρ is the persistence and σ_μ^2 the variance of the innovations.¹²

For the distortion, τ_{it} , we adopt a specification that allows for a rich correlation structure, both over time as well as with firm-level productivity. Specifically, τ_{it} takes the form

$$(6) \quad \tau_{it} = \gamma a_{it} + \varepsilon_{it} + \chi_i, \quad \varepsilon_{it} \sim \mathcal{N}(0, \sigma_\varepsilon^2), \quad \chi_i \sim \mathcal{N}(0, \sigma_\chi^2),$$

where the parameter γ controls the extent to which τ_{it} co-moves with productivity. If $\gamma < 0$, the distortion discourages (encourages) investment by firms with higher (lower) productivity, arguably, the empirically relevant case. The opposite is true if $\gamma > 0$. The uncorrelated component of τ_{it} has both permanent and i.i.d. (over time) components, denoted χ_i and ε_{it} , respectively. Thus, the severity of these factors is summarized by three parameters: $(\gamma, \sigma_\varepsilon^2, \sigma_\chi^2)$.¹³

Information.—Next, we spell out \mathcal{I}_{it} , the information set of the firm at the time of choosing K_{it+1} . This includes the entire history of productivity realizations through period t , i.e., $\{a_{it-s}\}_{s=0}^\infty$. Given the AR(1) assumption, this can be summarized by the most recent observation, namely, a_{it} . The firm also observes a noisy signal of the following period's innovation:

$$s_{it+1} = \mu_{it+1} + e_{it+1}, \quad e_{it+1} \sim \mathcal{N}(0, \sigma_e^2),$$

¹⁰The results in Section VA, where we solve a version of the model with nonconvexities without linearization, suggest that the perturbation yields reasonably accurate estimates.

¹¹We use lowercase to denote natural logs, a convention we follow throughout, so that, e.g., $x_{it} = \log X_{it}$.

¹²Online Appendix I.2 extends the analysis to allow for firm fixed effects in the process for a_{it} . This has no effect on our analytical results in Section II (where we work exclusively with growth rates) and Section IV. Our quantitative results regarding the sources of *arpk* dispersion are very similar with these effects.

¹³Online Appendix I.2 considers a more flexible process for τ_{it} . Our results there confirm the highly persistent nature of the uncorrelated component, suggesting that the simpler specification here with fixed and i.i.d. elements largely captures the time-series properties of the distortion.

where e_{it+1} is an i.i.d., mean-zero, and normally distributed noise term. This is in essence an idiosyncratic “news shock,” since it contains information about future productivity. Finally, firms also perfectly observe the uncorrelated transitory component of distortions, ε_{it+1} , at the time of choosing period t investment (as noted above, the distortion is indexed by the date it influences the firm’s capital choice, so that, e.g., ε_{it+1} is in the firm’s information set at date t and affects its choice of k_{it+1}). Firms also observe the fixed component of the distortion, χ_i . They do not see the correlated component, but are aware of its structure, i.e., they know γ .

Thus, the firm’s information set is given by $\mathcal{I}_{it} = (a_{it}, s_{it+1}, \varepsilon_{it+1}, \chi_i)$. Direct application of Bayes’ rule yields the conditional expectation of productivity, a_{it+1} :

$$a_{it+1} | \mathcal{I}_{it} \sim N(E_{it}[a_{it+1}], V)$$

where

$$E_{it}[a_{it+1}] = \rho a_{it} + \frac{V}{\sigma_e^2} s_{it+1}, \quad V = \left(\frac{1}{\sigma_\mu^2} + \frac{1}{\sigma_e^2} \right)^{-1}.$$

There is a one-to-one mapping between the posterior variance V and the noisiness of the signal, σ_e^2 (given the volatility of productivity, σ_μ^2). In the absence of “news,” i.e., $\sigma_e^2 = \infty$, we have $V = \sigma_\mu^2$, that is, posterior uncertainty is simply the variance of the innovation. This corresponds to a standard one period time-to-build assumption with $E_{it}[a_{it+1}] = \rho a_{it}$. At the other extreme, $\sigma_e^2 = 0$ implies $V = 0$, so the firm is perfectly informed about a_{it+1} . It turns out to be more convenient to work directly with the posterior variance, V , as the measure of uncertainty.

Law of Motion.—In online Appendix A.1, we solve the Euler equation in (4) to obtain

$$(7) \quad k_{it+1} = \psi_1 k_{it} + \psi_2 (1 + \gamma) E_{it}[a_{it+1}] + \psi_3 \varepsilon_{it+1} + \psi_4 \chi_i,$$

where

$$(8) \quad \xi(\beta\psi_1^2 + 1) = \psi_1((1 + \beta)\xi + 1 - \alpha),$$

$$\psi_2 = \frac{\psi_1}{\xi(1 - \beta\rho\psi_1)}, \quad \psi_3 = \frac{\psi_1}{\xi}, \quad \psi_4 = \frac{1 - \psi_1}{1 - \alpha}.$$

The coefficients ψ_1 – ψ_4 depend only on production (and preference) parameters, including the adjustment cost, and are independent of assumptions about information and distortions. The coefficient ψ_1 is increasing and ψ_2 – ψ_4 decreasing in the severity of adjustment costs, ξ . If there are no adjustment costs (i.e., $\xi = 0$), $\psi_1 = 0$ and $\psi_2 = \psi_3 = \psi_4 = \frac{1}{1 - \alpha}$. At the other extreme, as ξ tends to infinity, $\psi_1 \rightarrow 1$ and ψ_2 – ψ_4 vanish. Intuitively, as adjustment costs become large, the firm’s choice of capital becomes more autocorrelated and less responsive to productivity and distortions. Our empirical strategy essentially identifies the coefficients in the policy function, ψ_1 and $\psi_2(1 + \gamma)$, from observable moments. Given values of α and

β , the estimate of ψ_1 pins down ξ from (8). Given ξ, β , and ρ , we can use the estimate of $\psi_2(1 + \gamma)$ to recover γ .

Aggregation.—In online Appendix A.2, we show that aggregate output can be expressed as

$$\log Y \equiv y = a + \hat{\alpha}_1 k + \hat{\alpha}_2 n,$$

where k and n denote the (logs of the) aggregate capital stock and labor inputs, respectively. Aggregate TFP, denoted by a , is given by

$$(9) \quad a = a^* - \frac{(\theta \hat{\alpha}_1 + \hat{\alpha}_2) \hat{\alpha}_1}{2} \sigma_{arpk}^2, \quad \frac{da}{d\sigma_{arpk}^2} = -\frac{(\theta \hat{\alpha}_1 + \hat{\alpha}_2) \hat{\alpha}_1}{2},$$

where a^* is aggregate TFP if static capital products ($arpk_{it}$) are equalized across firms and σ_{arpk}^2 is the cross-sectional dispersion in (the log of) the static average product of capital ($arpk_{it} = p_{it}y_{it} - k_{it}$). Thus, aggregate TFP monotonically decreases in the extent of dispersion in capital productivities, summarized in this log-normal world by σ_{arpk}^2 . The effect of σ_{arpk}^2 on aggregate TFP depends on the elasticity of substitution, θ , and the relative shares of capital and labor in production. The higher is θ , that is, the closer we are to perfect substitutability, the more severe the losses from dispersion in capital products. Similarly, fixing the degree of overall returns to scale in production, for a larger capital share, $\hat{\alpha}_1$, a given degree of dispersion has larger effects on aggregate outcomes.¹⁴

In our framework, a number of forces (adjustment costs, information frictions, and distortions) will lead to $arpk$ dispersion. Once we quantify their contributions to σ_{arpk}^2 , equation (9) allows us to directly map those contributions to their aggregate implications.

Measuring the contribution of each factor is a challenging task, since all the data moments confound all the factors (i.e., each moment reflects the influence of more than one factor). As a result, there is no one-to-one mapping between individual moments and parameters: to accurately identify the contribution of any factor, we need to explicitly control for the others. In the following section, we overcome this challenge by exploiting the fact that these forces have different implications for different moments.

II. Identification

In this section, we lay out our identification strategy. Specifically, we provide a methodology to tease out the role of adjustment costs, informational frictions, and other factors using observable moments of firm-level value-added and investment. We use a tractable special case, when productivity follows a random walk, i.e., $\rho = 1$, to derive analytic expressions for key moments, allowing us to prove our identification result formally and make clear the underlying intuition. When we

¹⁴ Aggregate output effects are larger than TFP losses by a factor $\frac{1}{1 - \hat{\alpha}_1}$. This is because dispersion in capital products also reduces the incentives for capital accumulation and therefore, the steady-state capital stock.

return to the more general model (with $\rho < 1$) in the following section, we will demonstrate numerically that this intuition applies there as well.

We assume that the preference and technology parameters (the discount factor, β , the curvature of the profit function, α , and the depreciation rate, δ) are known to the econometrician (e.g., calibrated using aggregate or industry-level data). The remaining parameters of interest are the costs of capital adjustment, ξ , the quality of firm-level information (summarized by V), and the severity of distortions, parameterized by γ , σ_ε^2 , and σ_χ^2 .

Our methodology uses a set of carefully chosen elements from the covariance matrix of firm-level capital and productivity (the latter can be measured using data on value-added and capital, along with an assumed curvature, α). Note that $\rho = 1$ implies non-stationarity in levels, so we work with moments of (log) changes. This means that we cannot identify σ_χ^2 , the variance of the fixed component.¹⁵ Here, we focus on the four remaining parameters, namely ξ , γ , V , and σ_ε^2 . Our main result is to show that these are exactly identified by the following four moments: (i) the autocorrelation of investment, denoted $\rho_{k,k-1}$; (ii) the variance of investment, σ_k^2 ; (iii) the correlation of period t investment with the innovation in productivity in period $t - 1$, denoted $\rho_{k,a-1}$; and (iv) the coefficient from a regression of $\Delta arp_{k,it}$ on Δa_{it} , denoted $\lambda_{arpk,a}$.

Several of these moments have been used in the literature to quantify the various factors in isolation. For example, $\rho_{k,k-1}$ and σ_k^2 are standard targets in the literature on adjustment costs: see, e.g., Cooper and Haltiwanger (2006) and Asker, Collard-Wexler, and De Loecker (2014). The responsiveness to lagged fundamentals, $\rho_{k,a-1}$, is used by Klenow and Willis (2007) to quantify information frictions in a price-setting context. The covariance of $arpk$ with productivity, which we proxy with $\lambda_{arpk,a}$, is highlighted in the misallocation literature as suggestive of correlated distortions, e.g., Bartelsman, Haltiwanger, and Scarpetta (2013) and Buera and Fattal-Jaef (2018). The tractable random walk special case will shed light on the value of analyzing these moments/factors in tandem (and the potential biases from doing so in isolation).

Our main result is stated formally in the following proposition.

PROPOSITION 1: *The parameters ξ , γ , V , and σ_ε^2 are uniquely identified by the moments $\rho_{k,k-1}$, σ_k^2 , $\rho_{k,a-1}$, and $\lambda_{arpk,a}$.*

A. Intuition

The proof of Proposition 1 (in online Appendix A.3) involves tedious, if straightforward, algebra. Here, we provide a more heuristic argument for the intuition behind the result. We do this by analyzing parameters in pairs and showing that they can be uniquely identified by a pair of moments, holding the other parameters fixed. To be clear, this is a local argument; our goal here is simply to provide intuition about how the different moments can be combined to disentangle the different

¹⁵For our numerical analysis in Section III, we use a stationary model (i.e., with $\rho < 1$) and use σ_{arpk}^2 , a moment computed using levels of capital and productivity, to pin down σ_χ^2 .

forces. The identification result in Proposition 1 is a global one and shows that there is a unique mapping from the four moments to the four parameters.

Adjustment Costs and Correlated Distortions.—We begin with adjustment costs, parameterized by ξ , and correlated distortions, γ . The relevant moment pair is the variance and autocorrelation of investment, σ_k^2 and $\rho_{k,k-1}$. Both of these moments are commonly used to estimate quadratic adjustment costs, for example, Asker, Collard-Wexler, and De Loecker (2014) target the former and Cooper and Haltiwanger (2006) (among other moments) the latter. In our setting, these moments are given by

$$(10) \quad \sigma_k^2 = \left(\frac{\psi_2^2}{1 - \psi_1^2} \right) (1 + \gamma)^2 \sigma_\mu^2 + \frac{2\psi_3^2}{1 + \psi_1} \sigma_\varepsilon^2,$$


$$(11) \quad \rho_{k,k-1} = \psi_1 - \psi_3^2 \frac{\sigma_\varepsilon^2}{\sigma_k^2},$$

where $\psi_1 - \psi_3$ are as defined in equation (8). Our argument exploits the fact that the two forces have similar effects on the variability of investment, but opposing effects on the autocorrelation. To see this, recall that ψ_1 is increasing and ψ_2 and ψ_3 decreasing in adjustment costs, but all three are independent of γ . Thus, holding all other parameters fixed, σ_k^2 is decreasing in both the severity of adjustment costs (higher ξ) and correlated factors (more negative γ).¹⁶ The autocorrelation, $\rho_{k,k-1}$, on the other hand, increases with ξ but decreases as γ becomes more negative (through its effect on σ_k^2). Intuitively, while both forces dampen the volatility of investment, they do so for different reasons: adjustment costs make it optimal to smooth investment over time (increasing its autocorrelation) while correlated factors reduce sensitivity to the serially correlated productivity process (reducing the autocorrelation of investment).


Panel A of Figure 1 shows how these properties help identify the two parameters. The panel plots a pair of “isomoment” curves: each curve traces out combinations of the two parameters that give rise to a given value of the relevant moment, holding the other parameters fixed. Take the σ_k^2 curve: it slopes upward because higher ξ and lower γ have similar effects on σ_k^2 ; if γ is relatively small (in absolute value), adjustment costs must be high in order to maintain a given level of σ_k^2 . Conversely, a low ξ is consistent with a given value of σ_k^2 only if γ is very negative. An analogous argument applies to the $\rho_{k,k-1}$ isomoment curve: since higher ξ and more negative γ have opposite effects on $\rho_{k,k-1}$, the curve slopes downward. As a result, the two curves cross only once, yielding the unique combination of the parameters that is consistent with both moments. By plotting curves corresponding to the empirical values of these moments, we can uniquely pin down the pair (ξ, γ) (holding all other parameters fixed).

The graph also illustrates the potential bias introduced when examining these forces in isolation. For example, estimating adjustment costs while ignoring correlated distortions (i.e., imposing $\gamma = 0$) puts the estimate on the very right-hand


¹⁶The latter is true only for $\gamma > -1$, which is the empirically relevant region.

Panel A. Adjustment costs
versus correlated distortions

Panel B. Uncertainty versus correlated distortions



Panel C. Transitory versus correlated distortions



Panel D. Uncertainty versus adjustment costs




FIGURE 1. PAIRWISE IDENTIFICATION: ISOMOMENT CURVES

side of the horizontal axis. The estimate for ξ can be read off the vertical height of the isomoment curve corresponding to the targeted moment. Because the σ_k^2 curve is upward sloping, targeting this moment alone leads to an overestimate of adjustment costs (at the very right of the horizontal axis, the curve is above the point of intersection, which corresponds to the true value of the parameters).¹⁷ Targeting $\rho_{k,k-1}$ alone leads to a bias in the opposite direction: since the $\rho_{k,k-1}$ curve is downward sloping, imposing $\gamma = 0$ yields an underestimate of adjustment costs.

The remaining panels in Figure 1 repeat this analysis for other combinations of parameters. Each relies on the same logic as shown in panel A.

Uncertainty and Correlated Distortions.—To disentangle information frictions from correlated factors (panel B), we use the correlation of investment with past innovations in productivity, $\rho_{k,a-1}$, and the regression coefficient $\lambda_{arpk,a}$. These moments can be written as

$$(12) \quad \rho_{k,a-1} = \left[\frac{V}{\sigma_\mu^2}(1 - \psi_1) + \psi_1 \right] \frac{\sigma_\mu \psi_2 (1 + \gamma)}{\sigma_k},$$

$$(13) \quad \lambda_{arpk,a} = 1 - (1 - \alpha)(1 + \gamma)\psi_2 \left(1 - \frac{V}{\sigma_\mu^2} \right).$$

¹⁷This approach would also predict a counterfactually high level of the autocorrelation of investment.

A higher V implies a higher correlation of investment with lagged productivity innovations. Intuitively, the more uncertain is the firm, the greater the tendency for its actions to reflect productivity with a one-period lag. In contrast, a higher (more negative) γ increases the relative importance of transitory factors in the firm's investment decision, reducing its correlation with productivity. Therefore, to maintain a given level of $\rho_{k,a-1}$, a decrease in V must be accompanied by a less negative γ , i.e., the isomoment curve slopes downward. On the other hand, higher uncertainty and a more negative gamma both cause $arpk$ to covary more positively with contemporaneous productivity, a , leading to an upward sloping $\lambda_{arpk,a}$ curve. Together, these two curves pin down V and γ , holding other parameters fixed.

As before, the graph also reveals the direction of bias when estimating these factors in isolation. Assuming full information ($V = 0$) and using $\lambda_{arpk,a}$ to discipline the strength of correlated distortions, for example, as in Bartelsman, Haltiwanger, and Scarpetta (2013) and Buera and Fattal-Jaef (2018), overstates their importance. Using the lagged responsiveness to productivity to discipline information frictions while abstracting from correlated factors understates uncertainty.

Transitory and Correlated Distortions.—To disentangle correlated and uncorrelated factors, consider $\lambda_{arpk,a}$ and $\rho_{k,k-1}$. The former is increasing in the severity of correlated distortions, but independent of transitory ones, implying a vertical isomoment curve. The latter is decreasing in both types of distortions: a more negative γ and higher σ_ε^2 both increase the importance of the transitory determinants of investment, yielding an upward sloping isomoment curve.

Uncertainty and Adjustment Costs.—Finally, panel D shows the intuition for disentangling uncertainty from adjustment costs. An increase in the severity of either of these factors contributes to sluggishness in the response of actions to productivity, i.e., raises the correlation of investment with past productivity shocks $\rho_{k,a-1}$. However, the autocorrelation of investment $\rho_{k,k-1}$ is independent of uncertainty and determined only by adjustment costs (and other factors). Thus, holding those other factors fixed, the two moments, $\rho_{k,a-1}$ and $\rho_{k,k-1}$, jointly pin down the magnitude of adjustment frictions and the extent of uncertainty.

III. Quantitative Analysis

The analytical results in the previous section showed a tight relationship between the moments $(\rho_{k,a-1}, \rho_{k,k-1}, \sigma_k^2, \lambda_{arpk,a})$ and the parameters $(V, \xi, \sigma_\varepsilon^2, \gamma)$ for the special case of $\rho = 1$. In this section, we use this insight to develop an empirical strategy for the more general case where productivity follows a stationary AR(1) process and apply it to data on Chinese manufacturing firms. This allows us to quantify the severity of the various forces and their impact on $arpk$ dispersion and economic aggregates. For purposes of comparison, we also provide results for publicly traded firms in the United States.¹⁸ In Section IV, we extend our methodology to explore

¹⁸The two sets of firms are not directly comparable due to their differing coverage. For example, the Chinese data include many more small firms. Similarly, there may be selection biases when using data on publicly traded firms. To partly address this concern, in online Appendix J, we repeat the analysis on Chinese publicly traded

some specific candidates for firm-specific factors other than adjustment/informational frictions.

A. Parameterization

We begin by assigning values to the more standard preference and production parameters of our model. We assume a period length of one year and accordingly set the discount factor $\beta = 0.95$. We use an annual depreciation rate of $\delta = 0.10$. We keep the elasticity of substitution θ common across countries and set its value to 6, roughly in the middle of the range of values in the literature.¹⁹ We assume constant returns to scale in production, but allow the parameters $\hat{\alpha}_1$ and $\hat{\alpha}_2$ to vary across countries. In the United States, we set these to standard values of 0.33 and 0.67, respectively, which implies $\alpha = 0.62$.²⁰ For China, we set a higher capital share, namely, $\hat{\alpha}_1 = \hat{\alpha}_2 = 0.5$, in line with evidence from a number of recent papers, for example, Bai, Hsieh, and Qian (2006). These values imply an α equal to 0.71 in China.²¹

Next, we turn to the parameters of the productivity process, a_{it} : the persistence, ρ , and the variance of the innovations, σ_μ^2 . Under our assumptions, firm-level productivity is directly given by (up to an additive constant) $a_{it} = va_{it} - \alpha k_{it}$ where va_{it} denotes the log of value-added.²² Controlling for industry-year fixed effects to isolate the firm-specific component, we use a standard autoregression to estimate the parameters ρ and σ_μ^2 .

To pin down the remaining parameters (the adjustment cost, ξ , the quality of information, V , and the size of other factors, γ and σ_ε^2) we follow a strategy informed by the results in the previous section. Specifically, we target the correlation of investment growth with lagged innovations in productivity ($\rho_{i,a_{-1}}$), the autocorrelation of investment growth ($\rho_{i,i-1}$), the variance of investment growth (σ_i^2) and the correlation of the average product of capital with productivity ($\rho_{arpk,a}$).²³ Finally, to infer σ_χ^2 , the variance of the fixed component in (6), we match the overall dispersion in the average product of capital, σ_{arpk}^2 , which is clearly increasing in σ_χ^2 .

firms. We find patterns that are quite similar to those for Chinese manufacturing firms, suggesting that cross-country differences in frictions and distortionary factors are quite significant. This conclusion is further supported by results for two additional countries, Colombia and Mexico, also presented in online Appendix J.

¹⁹ In online Appendix I.3, we report results for $\theta = 3$, the value used in Hsieh and Klenow (2009).

²⁰ This is very close to the estimate of 0.59 in Cooper and Haltiwanger (2006). We also estimated α following the indirect inference approach in, e.g., Cooper, Gong, and Yan (2015). Specifically, we find the value of α so that the coefficient from an ordinary least squares (OLS) regression of value-added on capital using model-simulated data matches its counterpart from an identical regression in the data. This procedure also yields $\alpha = 0.62$.

²¹ Using the same capital share for both countries yields a very similar decomposition of observed σ_{arpk}^2 . More generally, the curvature of the profit function, α , plays a key role in determining the TFP/output implications of a given degree of σ_{arpk}^2 , but does not materially change the estimated contributions of various factors, the main focus of this paper. See also Section VB (where labor distortions leads to a higher α), online Appendix I.3 (where a lower elasticity of substitution leads to a lower α), as well as Section VD (sectoral heterogeneity in α).

²² An alternative strategy is to measure the true productivity directly, i.e., $\hat{a}_{it} = va_{it} - \alpha_1 k_{it} - \alpha_2 n_{it}$, and construct the implied $a_{it} = \frac{1}{1-\alpha_2} \hat{a}_{it}$. The two approaches are equivalent under an undistorted labor choice, but online Appendix E.1 shows that more generally, firm-specific capital profitability, a_{it} , is a combination of productivity and a labor distortion. As a result, inferring a_{it} from \hat{a}_{it} without adjusting for a potentially distorted labor choice can lead to biased estimates, while the strategy of directly measuring a_{it} , as we do here, remains valid.

²³ We use investment growth to partly cleanse the data of firm-level fixed effects, which have been shown to play a significant role in firm-level investment data (in the analytical cases studied earlier, we used the level of investment). See Morck, Shleifer, and Vishny (1990) for more on this issue. Online Appendix I.4 shows that our results are largely unchanged if we use the autocorrelation and variance of investment in levels, rather than growth rates.

TABLE 1—PARAMETERIZATION: SUMMARY

Parameter	Description	Target/value
Preferences/production		
θ	Elasticity of substitution	6
β	Discount rate	0.95
δ	Depreciation	0.10
$\hat{\alpha}_1$	Capital share	0.33 US/0.50 China
$\hat{\alpha}_2$	Labor share	0.67 US/0.50 China
Productivity/frictions		
ρ	Persistence of productivity	$\rho_{a,a-1}$
σ_μ^2	Shocks to productivity	
V	Signal precision	$\rho_{t,a-1}$
ξ	Adjustment costs	
γ	Correlated factors	
σ_ε^2	Transitory factors	
σ_χ^2	Permanent factors	σ_{ark}^2

Thus, by construction, our parameterized model will match the observed *arpk* dispersion in the data, allowing us to decompose the contribution of each factor. Online Appendix C describes our numerical estimation procedure in detail. We summarize our empirical approach in Table 1.

B. Data

The data on Chinese manufacturing firms are from the Annual Surveys of Industrial Production conducted by the National Bureau of Statistics. The surveys include all industrial firms (both state-owned and non-state-owned) with sales above 5 million RMB (about \$600,000). We use data spanning the period 1998–2009.²⁴ The original data come as a repeated cross section. A panel is constructed following almost directly the method outlined in Brandt, Van Biesebroeck, and Zhang (2014), which also contains an excellent overview of the data for the interested reader. The Chinese data have been used multiple times and are by now familiar in the misallocation literature (for example, Hsieh and Klenow 2009) although our use of the panel dimension is rather new. The data on United States publicly traded firms come from Compustat North America. We use data covering the same period as for the Chinese firms.

We measure the firm's capital stock, k_{it} , in each period as the value of fixed assets in China and of property, plant and equipment (PP&E) in the United States.²⁵ Value-added is estimated as a constant fraction of revenues using a share of intermediates of 0.5. We measure the average product of capital as $arpk_{it} = va_{it} - k_{it}$.

²⁴Industrial firms correspond to Chinese Industrial Classification codes 0610–1220, 1311–4392, and 4411–4620, which includes mining, manufacturing, and utilities. Early vintages of the NBS data did not report all variables for the full set of firms in the years after 2007. Although this does not seem to be an issue in our sample (all the variables we use are well populated in all years), we have also redone our analysis using data only through 2007. The estimates are very similar (as noted below, the moments are fairly stable over time).

²⁵Our baseline measure of the capital stock uses reported book values. In Section VD (details in online Appendix I.5), we construct the capital stock using the perpetual inventory method for the US firms and reestimate the model. This yields slightly different point estimates, but very similar patterns for the role of various factors.

TABLE 2—TARGET MOMENTS

	ρ	σ_μ^2	$\rho_{i,a_{-1}}$	$\rho_{i,l_{-1}}$	$\rho_{arpk,a}$	σ_i^2	σ_{arpk}^2
China	0.91	0.15	0.29	−0.36	0.76	0.14	0.92
United States	0.93	0.08	0.13	−0.30	0.55	0.06	0.45

Net investment and productivity growth are obtained by first differencing k_{it} and a_{it} , respectively. To isolate the firm-specific variation in our data series, we extract a time-by-industry fixed effect from each and use the residual. In both countries, industries are classified at the four-digit level. This is equivalent to deviating each firm from the unweighted average within its industry in each period and also eliminates aggregate components. After eliminating duplicates and problematic observations (for example, firms reporting in foreign currencies), outliers, observations with missing data, etc., our final sample consists of 797,047 firm-year observations in China and 34,260 in the United States. Online Appendix B provides further details on how we build our sample and construct the moments, as well as summary statistics from 2009.²⁶

Table 2 reports the target moments for both countries.²⁷ The first two columns show the productivity moments, which have similar persistence but higher volatility in China. The remaining columns show that, in China, investment growth is more correlated with past shocks, more volatile and less autocorrelated. The Chinese data also show a higher correlation between productivity and $arpk$ and substantially larger dispersion in $arpk$. These patterns will lead us to significantly different estimates of the severity of various factors across the two sets of firms.


C. Identification

Before turning to the estimation results, we revisit the issue of identification. Although we no longer have analytical expressions for the mapping between moments and parameters, we use a numerical experiment to show that the intuition developed in Section II for the random walk case applies here as well. In that section, we used a pairwise analysis to demonstrate how various moments combine to help disentangle the sources of observed $arpk$ dispersion. Here, we repeat that analysis by plotting numeric isomoment curves in Figure 2, using the moments and parameter values for US firms (from Tables 2 and 3, respectively). They reveal the same patterns as Figure 1, indicating that the logic of that special case goes through here as well.²⁸


²⁶ We have also examined the moments year-by-year. They are reasonably stable over time.

²⁷ We report bootstrapped standard errors for the moments in online Appendix Table C.1. Given the large sample sizes (almost 800,000 in China and 35,000 in the United States), the estimates of the moments are extremely precise.


²⁸ The differences in the precise shape of some of the curves in the two figures come partly from the departure from the random walk case and also from the fact that they use slightly different moments (Figure 2 works with changes in investment and $\rho_{arpk,a}$ while Figure 1 used changes in k and $\lambda_{arpk,a}$).

Panel A. Adjustment costs
versus correlated distortions

Panel B. Uncertainty versus correlated distortions



Panel C. Transitory versus correlated distortions



Panel D. Uncertainty versus adjustment costs




FIGURE 2. ISOMOMENT CURVES: QUANTITATIVE MODEL

D. The Sources of Misallocation

Table 3 contains our baseline results. In the top panel we display the parameter estimates.²⁹ In the second panel of Table 3, we report the contribution of each factor to *arpk* dispersion, which we denote $\Delta\sigma_{arpk}^2$.³⁰ These are calculated under the assumption that only the factor of interest is operational, i.e., in the absence of the others, so that the contribution of each one is measured relative to the undistorted first-best.³¹ The third panel expresses this contribution as a percentage of the total *arpk* dispersion measured in the data, denoted $\Delta\sigma_{arpk}^2/\sigma_{arpk}^2$. Because of interactions between the factors, there is no a priori reason to expect these relative contributions to sum to 1. In practice, however, we find that the total is reasonably close to 1, allowing us to interpret this exercise as a decomposition of total observed dispersion. In the bottom panel of the table, we compute the implied losses in aggregate

²⁹ Online Appendix Table C.1 reports standard errors and compares the model-simulated moments (at the estimated parameters) to their empirical counterparts. The parameters are quite precisely estimated (again, both of our firm-level datasets have a relatively large number of observations) and the model matches the five moments almost exactly in both countries.

³⁰ For adjustment costs, we do not have an analytic mapping between the severity of these costs and σ_{arpk}^2 , but this is a straightforward calculation to make numerically; for each of the other factors, we can compute their contributions to *arpk* dispersion analytically.

³¹ An alternative would be to calculate the contribution of each factor holding the others constant at their estimated values. It turns out that the interactions between the factors are small at the estimated parameter values, so the two approaches yield similar results. Online Appendix Table D.1 shows that the effects of each factor on *arpk* dispersion in the United States are close under either approach. Interaction effects are even smaller in China.

TABLE 3—CONTRIBUTIONS TO “MISALLOCATION”

	Adjustment costs	Uncertainty	Other factors		
			Correlated	Transitory	Permanent
<i>Parameters</i>	ξ	V	γ	σ_{ε}^2	σ_{χ}^2
China	0.13	0.10	−0.70	0.00	0.41
United States	1.38	0.03	−0.33	0.03	0.29
$\Delta\sigma_{arpk}^2$					
China	0.01	0.10	0.44	0.00	0.41
United States	0.05	0.03	0.06	0.03	0.29
$\Delta\sigma_{arpk}^2/\sigma_{arpk}^2$					
China	1.3%	10.3%	47.4%	0.0%	44.4%
United States	10.8%	7.3%	14.4%	6.3%	64.7%
Δa					
China	0.01	0.08	0.38	0.00	0.36
United States	0.02	0.01	0.03	0.01	0.13

TFP, again relative to the undistorted first-best level, i.e., $\Delta a = a^* - a$. Once we have the contribution of each factor to *arpk* dispersion, computing these values is simply an application of expression (9).

Adjustment Costs.—Our results show evidence of economically significant adjustment frictions. For example, the estimate of 1.38 for ξ in the United States implies a value of 0.2 for $\hat{\xi}$ in the adjustment cost function.³² This puts us in the middle of previous estimates of convex costs in the literature, though differences in data and empirical strategies complicate direct comparisons. For example, using US manufacturing data, Asker, Collard-Wexler, and De Loecker (2014) estimate a convex adjustment cost of 8.8 in a monthly model, which translates to an annual $\hat{\xi} = 0.73$, roughly a factor of four above our estimate.³³ Our estimate is closer to, and slightly higher than, Cooper and Haltiwanger (2006), who find $\hat{\xi} = 0.05$ for US manufacturing plants.

What leads us to find different estimates? The answer lies primarily in the fact that our model explicitly includes additional factors that may act on the investment decision (e.g., distortions) and consequently, our empirical strategy is designed to match a broader set of moments.³⁴ The papers mentioned above abstract from these factors and focus on matching different moments. For example, Asker, Collard-Wexler, and De Loecker (2014) target the variability of investment (among

³²The mapping between ξ and $\hat{\xi}$ is in equation (A.1) in online Appendix A.1.

³³To interpret this difference, a firm that doubles its capital stock in a year would incur an adjustment cost equal to 11 percent of the value of the investment according to our estimate, but equal to 60 percent at the Asker, Collard-Wexler, and De Loecker (2014) estimate.

³⁴There are a few other differences between our approach and these papers: (i) they have convex and nonconvex (fixed) adjustment costs. In Section VA, we show that our estimates of $\hat{\xi}$ change little when we introduce a fixed cost; (ii) they use moments of investment in levels while we work with growth rates. In online Appendix I.4, we show that targeting the variance and autocorrelation of investment in levels changes the estimate of $\hat{\xi}$ only slightly; (iii) Asker, Collard-Wexler, and De Loecker (2014) follow a different strategy to estimate the process for profitability, a_{it} : they directly measure productivity \hat{a}_{it} and use the implied a_{it} . See footnote 22. As we show in online Appendix E.1, this strategy can overstate the volatility of a_{it} , i.e., σ_{μ}^2 , and bias adjustment cost estimates upward.

other moments), but do not try to match the autocorrelation, while Cooper and Haltiwanger (2006) do the reverse. As we saw in Section II, in the presence of correlated factors, the first strategy overstates the true extent of adjustment costs, while the second understates it. It turns out that this bias can be quite large: an adjustment cost-only model (i.e., ignoring other factors) estimated to match the volatility of investment growth yields an estimate of $\hat{\xi}$ about 60 percent higher than our baseline estimate, but predicts a counterfactually high autocorrelation of investment growth: -0.17 versus -0.30 in the data. A strategy targeting only the serial correlation leads to the opposite conclusion: a lower estimate of $\hat{\xi}$, but at the cost of excessively high variability compared to the data. These patterns are exactly in line with the arguments developed in Section II. More broadly, these exercises can partly explain the wide range of adjustment cost estimates in the literature: when adjustment costs are estimated without explicitly controlling for other factors, the results can be quite sensitive to the particular moments chosen.³⁵ Indeed, our results suggest that explicitly accounting for these additional factors is essential in order to reconcile a broad set of moments in firm-level investment dynamics.

The estimated value of ξ is lower in China. Investment growth in China is both more volatile and less serially correlated than for US firms, which (together with the other moments) leads the estimation to find a lower degree of adjustment frictions. Importantly, as in the United States, one would reach a very different conclusion from examining a model with only adjustment costs: for example, estimating such a model by targeting σ_i^2 in China yields an estimate for ξ of about 1.5, roughly 10 times larger than the one in Table 3.

Perhaps most importantly for purposes of our analysis, in both countries, the estimated adjustment costs do not contribute significantly to *arpk* dispersion. If adjustment costs were the only friction in China, σ_{arpk}^2 would be 0.01 (the observed level is 0.92). The higher estimate of ξ in the United States implies a slightly higher, though still modest, contribution (by themselves, adjustment costs lead to $\sigma_{arpk}^2 = 0.05$ or 11 percent of the observed dispersion). The corresponding aggregate TFP losses are 1 and 2 percent in the two countries, respectively.

This does not mean that adjustment costs are irrelevant for understanding firm-level investment dynamics. Setting adjustment costs to zero in the United States while holding the other parameters at their estimated values causes the variance of investment growth to spike to 1.68 (compared to 0.06 in the data) and the autocorrelation to plummet to -0.62 (data: -0.30). However, σ_{arpk}^2 falls only modestly, from 0.45 to 0.41. Reestimating the model without adjustment costs (and dropping the autocorrelation as a target) also leads to a counterfactually low autocorrelation (-0.50).³⁶ In other words, while adjustment frictions are an important determinant of investment dynamics, they do not generate significant dispersion in average products of capital.³⁷

³⁵ Bloom (2009) points out the wide variation in these estimates, ranging from 0 to 20 (Table IV).

³⁶ The estimates for other parameters also change: notably, a more negative γ is needed to match σ_i^2 .

³⁷ Asker, Collard-Wexler, and De Loecker (2014) make a similar observation: across various specifications of adjustment costs (including one with zero adjustment costs and a one-period time-to-build), their model's performance in capturing dispersion in *arpk* is not dramatically altered, even though the implications for other moments (e.g., the variability of investment) are quite different. See Table 9 and the accompanying discussion in that paper.

Uncertainty.—Table 3 shows that firms in both countries make investment decisions under considerable uncertainty, with the information friction more severe for Chinese firms. As a share of the prior uncertainty, σ_μ^2 , residual uncertainty, V/σ_μ^2 , is 0.42 in the United States and 0.63 in China.³⁸ In an environment where imperfect information is the only friction, we have $\sigma_{arpk}^2 = V$, so the contribution of uncertainty alone to observed *arpk* dispersion can be directly read off the second column in Table 3, namely 0.10 in China and 0.03 in the United States. These represent about 10 and 7 percent of total *arpk* dispersion in the two countries, respectively. The implications for aggregate TFP are substantial in China (losses are about 8 percent) and are lower in the United States, about 1 percent. Note, however, that imposing a one-period time-to-build assumption where firms install capital in advance without any additional information about innovations in productivity, i.e., setting $V = \sigma_\mu^2$, would overstate uncertainty (and bias the estimates of adjustment costs and other parameters). Indeed, doing so yields estimates of V that are about 55 percent higher in China and a factor of 2.5 times higher in the United States.

Distortions.—The last three columns of Table 3 show that other, potentially distortionary, factors play a significant role in generating the observed *arpk* dispersion in both countries. Turning first to the correlated component, the negative values of γ suggest that they act to disincentivize investment by more productive firms and especially so in China. The contribution of these distortions to *arpk* dispersion is given by $\gamma^2 \sigma_a^2$, which amounts to 0.44 in China, or 47 percent of total dispersion. The associated aggregate consequences are also quite sizable: TFP losses from these sources are 38 percent. In contrast, the estimate of γ in the United States is significantly less negative than in China, suggesting that these types of correlated factors are less of an issue for firms in the United States, both in an absolute sense (the *arpk* dispersion from these factors in the United States is 0.06, less than one-seventh that in China) and in relative terms (they account for only 14 percent of total observed *arpk* dispersion in the United States). The corresponding TFP effects are also considerably smaller for the United States: losses from correlated sources are only about 3 percent.

Next, we consider the role of distortions that are uncorrelated with firm productivity. Table 3 shows that purely transitory factors (measured by σ_ε^2) are negligible in both countries, but permanent firm-specific factors (measured by σ_χ^2) play a prominent role. Their contribution to *arpk* dispersion, which is also given by σ_χ^2 , amounts to 0.41 in China and 0.29 in the United States. Thus, their absolute magnitude in the United States is considerably below that in China, but in relative terms, these factors seem to account for a substantial portion of measured *arpk* dispersion in both countries. The aggregate consequences of these types of distortions are also

³⁸ Our values for V/σ_μ^2 are similar to those in David, Hopenhayn, and Venkateswaran (2016), who find 0.41 and 0.63 for publicly traded firms in the United States and China, respectively. The estimates of V are different but are not directly comparable: David, Hopenhayn, and Venkateswaran (2016) focus on longer time horizons (they analyze 3-year time intervals). This might lead one to conclude that ignoring other factors, as David, Hopenhayn, and Venkateswaran (2016) do, leads to negligible bias in the estimate of uncertainty. But, this is not a general result and rests on the fact that adjustment costs and uncorrelated distortions are estimated to be modest. Then, as Figure 2 shows, the sensitivity of actions to signals turns out to be a very good indicator of uncertainty. If, on the other hand, adjustment costs and/or uncorrelated factors were much larger, the bias from estimating uncertainty alone can be quite significant.

significant, with TFP losses of 36 percent in China and about 13 percent in the United States.

In sum, the estimation results point to the presence of substantial distortions to investment, especially in China, where they disproportionately disincentivize investment by more productive firms. Section V and online Appendix I show that these results are robust to a number of modifications to our baseline setup, e.g., allowing for nonconvex adjustment costs, a frictional labor choice, richer stochastic processes on productivity and distortions, curvature assumptions and additive measurement error. Further, we have applied the methodology to data on Colombian and Mexican firms (in addition to the set of publicly traded firms in China): the results resemble those for Chinese manufacturing firms, in that they point to a substantial role for correlated factors, as well as fixed ones (details are in online Appendix J).

What patterns in the data lead us to this conclusion? In both countries, we see considerable dispersion in $arpk$, which tends to be correlated with firm productivity. The fact that investment growth is not very correlated through time and responds only modestly to past shocks limits the role of adjustment and informational frictions and assigns a substantial role to other factors. The high correlation of $arpk$ with productivity, particularly in China, suggests that these factors vary systematically with productivity. In both countries, large fixed distortions, uncorrelated with productivity, are necessary to rationalize the total observed dispersion in $arpk$. In the next section, we explore some candidates for these firm-specific factors.

IV. Firm-Specific Factors: Some Candidates

Our results suggest a large role for firm-specific “distortions” in explaining the observed $arpk$ dispersion. In this section, we extend our baseline framework and empirical methodology to investigate three potential sources: heterogeneity in markups and production technologies, size-dependent policies and financial considerations.

A. Heterogeneity in Markups and Technologies

In our baseline setup, all firms within an industry (i) operated identical production technologies and (ii) were monopolistically competitive facing CES demand curves and therefore, had identical markups. As a result, any firm-level heterogeneity in technologies and/or markups would show up in our estimates of other factors (note also that this type of variation would drive a wedge between $arpk$ and the true marginal product of capital, implying that dispersion in the former is not necessarily a sign of misallocation). Here, we explore this possibility using a modified version of our baseline model. This requires more assumptions and additional data, but allows us to provide an upper bound on the contribution of these elements.

We begin by generalizing the production function from Section I to include intermediate inputs and to allow for (potentially time-varying) heterogeneity in input intensities. Specifically, the output of firm i is now given by

$$Y_{it} = K_{it}^{\hat{\alpha}_{it}} N_{it}^{\hat{\zeta}_{it} - \hat{\alpha}_{it}} M_{it}^{1 - \hat{\zeta}_{it}},$$

where M_{it} denotes intermediate or materials input. The price of these inputs, potentially firm-specific, is denoted P_{it}^M . Throughout this section, we abstract from adjustment/information frictions in firms' input decisions. This is largely for simplicity, but is also supported by the relatively modest role played by these dynamic considerations in our baseline estimates.³⁹

Capital and labor choices are assumed to be subject to a factor-specific "distortion" (in addition to the markup), denoted T_{it}^K and T_{it}^N , respectively, but the choice of intermediates is undistorted except for the markup. Since the method remains valid even with unobserved firm-specific variation in the price of intermediate goods, it does allow for distortions in the market for intermediate inputs, so long as they are reflected in prices. Formally, these assumptions imply that the firm's optimal choices solve the following cost minimization problem:

$$\min_{K_{it}, N_{it}, M_{it}} R_t T_{it}^K K_{it} + W_t T_{it}^N N_{it} + P_{it}^M M_{it},$$

subject to

$$Y_{it} \leq K_{it}^{\hat{\alpha}_{it}} N_{it}^{\hat{\zeta}_{it} - \hat{\alpha}_{it}} M_{it}^{1 - \hat{\zeta}_{it}}.$$

The Contribution of Markup Dispersion.—To quantify the contribution of markup dispersion, we use the powerful methodology of De Loecker and Warzynski (2012), which allows us to measure firm-level markups without taking a stand on the nature of competition/demand. To do so, we assume a common materials elasticity across firms within an industry, i.e., $\hat{\zeta}_{it} = \hat{\zeta}_t \forall i$.⁴⁰ Cost minimization then implies the following optimality condition (we suppress the time subscript on ζ , though the method remains valid with arbitrary time-variation):

$$(14) \quad P_{it}^M = MC_{it} (1 - \hat{\zeta}) \frac{Y_{it}}{M_{it}} \Rightarrow \frac{P_{it}^M M_{it}}{P_{it} Y_{it}} = (1 - \hat{\zeta}) \frac{MC_{it}}{P_{it}},$$

where MC_{it} is the marginal cost of the firm. This condition states that, at the optimum, the firm sets the materials share of revenue equal to the inverse of the markup, MC_{it}/P_{it} , multiplied by the materials elasticity $1 - \hat{\zeta}$.

Expression (14) suggests a simple way to estimate the cross-sectional dispersion in markups. The left-hand side is materials' share of revenue: the within-industry dispersion in this object (in logs) maps one-for-one into (log) markup dispersion. Data on materials expenditures are directly reported in the Chinese data. In the United States, we follow, e.g., De Loecker and Eeckhout (2017) and İmrohoroglu and Tüzel (2014) and calculate intermediate expenditures as total expenses less

³⁹It is possible to extend the identification methodology from Section II to explicitly incorporate these forms of heterogeneity. Although this would require more assumptions (e.g., on the correlation structure of markups/technologies with productivity and over time) and make the intuition more complicated, the basic insights should still go through.

⁴⁰In online Appendix F, we allow the materials elasticity to vary across firms within an industry (our baseline calculation already allows for variation across industries). Under certain orthogonality assumptions, we show that the covariance of the materials share with *arpk* and *arpn* (average product of labor) pins down the dispersion in markups. This approach yields very similar (if slightly lower) estimates of markup dispersion in both countries.

labor expenses. The former are defined as sales less operating income and the latter are imputed using number of employees and the average industry wage, from the NBER-CES Manufacturing Industry Database.⁴¹

The results of applying this procedure are reported in Table 4. The variance of the share of materials in revenue (the first row in the top panel) is about 0.06 in the US Compustat data and 0.05 in China. This accounts for about 14 percent of σ_{arpk}^2 among the United States firms, but only about 4 percent among Chinese manufacturing firms. Thus, markup heterogeneity composes a nonnegligible fraction of observed *arpk* dispersion among United States publicly traded firms but seems to be an almost negligible force in China.

Can markup variation help explain the large role for correlated distortions in Section III? In theory, yes: markups that increase with size would be consistent with the patterns uncovered in that section. Quantitatively, however, this does not seem to hold much promise. In China, this is clear simply from the rather modest dispersion in markups: the contribution of markups to *arpk* dispersion is much smaller (0.05) than the estimated total contribution of correlated factors (0.44). In the United States, where dispersion in markups is more substantial, the data suggest they are largely independent of firm size. For example, projecting the measured markup on revenues yields a statistically significant, yet economically small, coefficient of 0.01 (the raw correlation between markups and revenues is 0.07).

The Contribution of Technology Dispersion.—Cost minimization also implies that the average revenue products of capital and labor are given by⁴²

$$(15) \quad arpk_{it} \equiv \log\left(\frac{P_{it}Y_{it}}{K_{it}}\right) = \log\frac{P_{it}}{MC_{it}} - \log\hat{\alpha}_{it} + \tau_{it}^K + \text{Constant},$$

$$(16) \quad arpn_{it} \equiv \log\left(\frac{P_{it}Y_{it}}{N_{it}}\right) = \log\frac{P_{it}}{MC_{it}} - \log(\hat{\zeta} - \hat{\alpha}_{it}) + \tau_{it}^N + \text{Constant}$$

$$(17) \quad \approx \log\frac{P_{it}}{MC_{it}} + \left(\frac{\bar{\alpha}}{\hat{\zeta} - \bar{\alpha}}\right)\log\hat{\alpha}_{it} + \tau_{it}^N + \text{Constant},$$

where τ_{it}^K and τ_{it}^N denote (the logs of) the capital and labor wedges T_{it}^K and T_{it}^N , respectively, and $\bar{\alpha}$ is the average capital elasticity. Average revenue products are combinations of firm-specific production elasticities, markups and wedges. Note that the capital elasticity, $\hat{\alpha}_{it}$, has opposing effects on the average products of capital and labor: firms with a high $\hat{\alpha}_{it}$ will, ceteris paribus, have a low *arpk* and a high *arnp*. In other words, this form of heterogeneity acts like a “mix” distortion (as opposed to “scale” factors, which distort all input decisions in the same direction). We will

⁴¹Details of these calculations are provided in online Appendix B. In an earlier version of this paper, we also used the reported wage bill, available for a smaller subset of firms. The results were broadly similar, though the share of both markup and technology dispersion was somewhat higher (28 and 62 percent of σ_{arpk}^2 for those firms).

⁴²See online Appendix F for details. The third equation is derived by log-linearizing (16) around $\hat{\alpha}_{it} = \bar{\alpha}$.

TABLE 4—HETEROGENEOUS MARKUPS AND TECHNOLOGIES

	China	United States		
<i>Moments</i>				
$\sigma^2\left(\log \frac{P_{it}Y_{it}}{P_{it}^M M_{it}}\right)$	0.05	0.06		
$\text{cov}\left(\overline{arpk_{it}}, \overline{arpn_{it}}\right)$	0.41	0.23		
$\sigma^2\left(\overline{arpk_{it}}\right)$	1.37	0.52		
$\sigma^2\left(\overline{arpn_{it}}\right)$	0.76	0.35		
<i>Estimated $\Delta\sigma_{arpk}^2$</i>				
Dispersion in markups	0.05	(3.8%)	0.06	(13.6%)
Dispersion in $\log \hat{\alpha}_{it}$	0.30	(23.1%)	0.18	(44.4%)
Total	0.35	(26.9%)	0.24	(58.1%)

Note: The values in parentheses in the bottom panel are the contributions to $arpk$ dispersion expressed as a fraction of total σ_{arpk}^2 .

make use of this property to derive an upper bound for variation in $\hat{\alpha}_{it}$ using the observed covariance of $arpk$ and $arpn$. Let

$$\widetilde{arpk_{it}} \equiv \log\left(\frac{P_{it}Y_{it}}{K_{it}}\right) - \log\left(\frac{P_{it}}{MC_{it}}\right),$$

$$\widetilde{arpn_{it}} \equiv \log\left(\frac{P_{it}Y_{it}}{N_{it}}\right) - \log\left(\frac{P_{it}}{MC_{it}}\right),$$

denote the markup-adjusted average revenue products of capital and labor. Online Appendix F proves the following result.

PROPOSITION 2: *Suppose $\log \hat{\alpha}_{it}$ is uncorrelated with the distortions τ_{it}^K and τ_{it}^N . Then, the cross-sectional dispersion in $\log \hat{\alpha}_{it}$ satisfies*

$$(18) \quad \sigma^2(\log \hat{\alpha}_{it}) \leq \frac{\sigma_{arpk}^2 \sigma_{arpn}^2 - \text{cov}(\widetilde{arpk}, \widetilde{arpn})^2}{2 \frac{\bar{\alpha}}{\hat{\zeta} - \bar{\alpha}} \text{cov}(\widetilde{arpk}, \widetilde{arpn}) + \left(\frac{\bar{\alpha}}{\hat{\zeta} - \bar{\alpha}}\right)^2 \sigma_{arpk}^2 + \sigma_{arpn}^2}.$$

The bound in (18) is obtained by setting the correlation between τ_{it}^K and τ_{it}^N to 1. Given the observed second moments of $(\widetilde{arpk_{it}}, \widetilde{arpn_{it}})$, this maximizes the potential for variation in $\hat{\alpha}_{it}$, which, as noted earlier, is a source of negative correlation between $\widetilde{arpk_{it}}$ and $\widetilde{arpn_{it}}$. The expression for the bound reveals the main insight: the more positive the covariance between $(\widetilde{arpk_{it}}, \widetilde{arpn_{it}})$, the lower is the scope for heterogeneity in $\hat{\alpha}_{it}$.

To compute this bound for the two countries, we set $\hat{\zeta}$, the share of materials in gross output, to 0.5. The results, along with the moments used, are reported in Table 4. They show that heterogeneity in technologies can potentially account for a substantial portion of σ_{arpk}^2 in the United States, as much as 44 percent, and a more modest, though still significant, fraction in China, about 23 percent.⁴³ In other words, a

⁴³There is some evidence that the share of intermediates may be higher in China than the United States, see, e.g., Table 1 in Brandt, Van Biesebroeck, and Zhang (2014). We recomputed the bound with $\hat{\zeta} = 0.25$ and obtained

substantial portion of measured *arpk* dispersion in both countries may not be a sign of misallocated resources at all. This is most striking in the United States, where the average products of capital and labor covary less positively than in China.⁴⁴

An alternative approach to assessing the potential for technology dispersion is discussed in Hsieh and Klenow (2009), in which all the variation in firm-level capital-labor ratios is attributed to heterogeneity in $\hat{\alpha}_{it}$. This amounts to assuming that $\tau_{it}^K = \tau_{it}^N$, which implies

$$\begin{aligned} k_{it} - n_{it} &= arpn_{it} - arpk_{it} \approx \frac{\hat{\zeta}}{\hat{\zeta} - \bar{\alpha}} \log \hat{\alpha}_{it} \Rightarrow \\ \sigma^2(k_{it} - n_{it}) &= \left(\frac{\hat{\zeta}}{\hat{\zeta} - \bar{\alpha}} \right)^2 \sigma^2(\log \hat{\alpha}_{it}). \end{aligned}$$

This procedure yields estimates for $\sigma^2(\log \hat{\alpha}_{it})$ that are quite close to those in Table 4: 0.27 (compared to 0.30) for China and 0.16 (compared to 0.18) in the United States.

In sum, unobserved heterogeneity in markups and technologies seem to be promising candidates for firm-specific factors, which drive most of the *arpk* dispersion in the data. This is particularly true for the United States, where they can explain as much as 58 percent of the observed dispersion. In China, their role is more modest, but still meaningful, at 27 percent.

B. Size-Dependent Policies

Our results show a significant role for factors correlated with firm-level productivity, especially in China, in explaining the observed variation in *arpk* across firms. Here, we show how policies that affect or restrict the size of firms can lead to a correlated factor of this form. A number of papers have pointed out the prevalence of distortionary size-dependent policies across a range of countries, for example, Guner, Ventura, and Xu (2008). These policies often take the form of restrictions (or additional costs) associated with acquiring capital and/or other inputs. To be clear, our goal is not to explore the role of a particular policy in China or the United States. Rather, we show how policies that are common in a number of countries can generate patterns that are, in a sense, isomorphic to factors correlated with productivity.

Toward this end, we generalize our baseline specification of firm-specific factors in equation (6) to allow for a component that varies with the chosen level of capital. Formally,

$$\tau_{it} = \gamma_k k_{it} + \gamma a_{it} + \varepsilon_{it} + \chi_i,$$

where the parameter γ_k indexes the severity of this additional component. The empirically relevant case is $\gamma_k < 0$, which implicitly penalizes larger firms. This

very similar results. We also verified the accuracy of the approximation by working directly with (16) instead of the log-linearized version in (17). This yielded slightly lower bounds: 38 and 17 percent of σ_{arpk}^2 in the United States and China, respectively.

⁴⁴In online Appendix F, we derive an analogous bound in the case with within-industry variation in the materials elasticity. The results are very similar.

TABLE 5—SIZE VERSUS PRODUCTIVITY-DEPENDENT FACTORS

	Correlated factors			Adj. costs ξ
	Size-dependent	Prod.-dependent	Total	
	γ_k	γ		
$\alpha + \gamma_k = 0.71$ (<i>baseline</i>)				
Parameters	0.00	−0.70		0.13
$\Delta\sigma_{arpk}^2/\sigma_{arpk}^2$	0.0%	47.4%	47.4%	1.3%
$\alpha + \gamma_k = 0.54$				
Parameters	−0.18	−0.51		0.21
$\Delta\sigma_{arpk}^2/\sigma_{arpk}^2$	14.2%	25.4%	39.6%	2.3%
$\alpha + \gamma_k = 0.36$				
Parameters	−0.36	−0.33		0.29
$\Delta\sigma_{arpk}^2/\sigma_{arpk}^2$	29.6%	10.2%	39.8%	3.2%

specification captures the essence of the policies discussed above in a tractable way (e.g., it allows us to continue to use perturbation methods). With this formulation, the log-linearized Euler equation takes the form

$$(19) \quad k_{it+1}((1 + \beta)\xi + 1 - \alpha - \gamma_k) \\ = (1 + \gamma)E_{it}[a_{it+1}] + \varepsilon_{it+1} + \chi_i + \beta\xi E_{it}[k_{it+2}] + \xi k_{it}.$$

Expression (19) is identical to expression (4), but with $\alpha + \gamma_k$ taking the place of α . It is straightforward to derive the firm's investment policy function and verify that the same adjustment goes through, i.e., expressions (7) and (8) hold, with α everywhere replaced by $\alpha + \gamma_k$. Intuitively, the size-dependent component, γ_k , changes the effective degree of curvature in the firm's investment problem; although the curvature of the profit function remains α , the firm acts as if it is $\alpha + \gamma_k$. If $\gamma_k < 0$, the distortion dampens the responsiveness of investment to shocks. If $\gamma_k > 0$, the responsiveness of investment is amplified.

Importantly, these effects are broadly similar to those coming from γ : indeed, if γ_k were the only factor distorting investment choices, the implied law of motion for k_{it} is identical (up to a first-order) to one with only productivity-dependent factors, where $\gamma = \frac{\gamma_k}{1 - \alpha - \gamma_k}$. The implication of this isomorphism is that we cannot distinguish the two factors using observed series of capital and value-added alone. This challenge also applies to the case when other factors are present, though the mapping between the two is more complicated (and affects the other parameters as well). We detail this mapping in online Appendix G.

What about the contribution to *arpk* dispersion? Table 5 reports the results for Chinese firms for two values of γ_k , namely −0.18 and −0.36 (these values imply effective curvatures $\alpha + \gamma_k$ equal to one-quarter and one-half of the true α , respectively). The table shows two key results: first, a more negative γ_k reduces the estimated γ (i.e., makes it less negative), suggesting that our baseline estimates of correlated factors could be picking up such size-dependent policies. The total contribution of both types of distortions is quite stable, ranging between 40 and

47 percent. Second, the estimates of adjustment costs remain modest over this wide range of curvature.

C. Financial Frictions

In this section, we show that liquidity considerations can lead to size-dependent distortions of the form analyzed in the previous subsection. We assume that firms face a cost $\Upsilon(K_{it+1}, B_{it+1})$, where B_{it+1} denotes holdings of liquid assets, which earn an exogenous rate of return $R < 1/\beta$. The cost is increasing (decreasing) in K_{it+1} (B_{it+1}). This specification captures the idea that firms need costly liquidity in order to operate (e.g., to meet working capital needs). Using a continuous penalty function rather than an occasionally binding constraint allows us to continue using perturbation methods. Note also that this differs from the standard borrowing constraint used widely in the literature on financial frictions. Our firms are not constrained in terms of their ability to raise funds. This implies that self-financing, which often significantly weakens the long-run bite of borrowing constraints, plays no role here.⁴⁵

We use the following flexible functional form for the liquidity cost:

$$\Upsilon(K_{it+1}, B_{it+1}) = \hat{\nu} \frac{K_{it+1}^{\omega_1}}{B_{it+1}^{\omega_2}},$$

where $\hat{\nu}$, ω_1 , and ω_2 are all positive parameters. The marginal liquidity cost of capital, after optimizing over the choice of B_{it+1} is given by (derivations in online Appendix H)

$$(20) \quad \Upsilon_{1,t+1} \equiv \frac{d\Upsilon(K_{it+1}, B_{it+1})}{dK_{it+1}} = \nu(1 - \beta R)^{\frac{\omega_2}{\omega_2+1}} K_{it+1}^{\omega},$$

where ν and ω are composite parameters. The former is always positive, while the latter is of indeterminate sign. If ω is positive (negative), the marginal cost of liquidity is increasing (decreasing) in K_{it+1} .

The log-linearized Euler equation takes the same form as (19), with

$$(21) \quad \gamma_k = -\omega \left(\frac{\bar{\Upsilon}_1}{\bar{\Upsilon}_1 + \kappa} \right),$$

where $\bar{\Upsilon}_1$ is the marginal cost of liquidity in the deterministic steady state and $\kappa = 1 - \beta(1 - \delta) + \hat{\xi}\delta\left(1 - \beta\left(1 - \frac{\delta}{2}\right)\right)$. Intuitively, the fraction $\frac{\bar{\Upsilon}_1}{\bar{\Upsilon}_1 + \kappa}$ is the steady state share of liquidity in the total marginal cost of capital. Thus, liquidity considerations manifest themselves as a size-dependent factor of the form described in Section IVB. The sign depends on the sign of ω : if $\omega > 0$, then $\gamma_k < 0$, so costly liquidity dampens incentives to adjust capital in response to productivity (since the liquidity cost is convex). The opposite happens if $\omega < 0$.

Thus, liquidity considerations are a promising candidate for correlated and/or size-dependent factors. Cross-country differences in liquidity requirements

⁴⁵ See, for example, Midrigan and Xu (2014) and Moll (2014). Gopinath et al. (2017) show that a richer variant of the standard collateral constraint can have important implications during a period of transition, even if it generates only modest amounts of *arpk* dispersion in the long-run.

(summarized by the parameters ν and ω) and/or costs (i.e., $1 - \beta R$) will translate into variation in the severity of our measures of correlated firm-specific factors. However, our results here also highlight the difficulty in separating them from other factors using production-side data alone. One would need additional data, e.g., on firm-level liquidity holdings, to disentangle the role of liquidity from other forces.

In sum, our findings in Sections IVA–IVC provide some guidance on the factors beyond adjustment and information frictions that influence investment decisions. For US publicly traded firms, observed dispersion in value-added/capital ratios could be driven to a significant extent by unobserved heterogeneity in production technologies and therefore, as we emphasized above, may not be a sign of misallocated capital. On the other hand, the scope for this type of heterogeneity appears limited among Chinese manufacturing firms, suggesting a greater role for inefficient factors like size-dependent policies or financial imperfections.

V. Robustness and Extensions

In this section, we explore a number of variants on our baseline approach. We generalize our specification of adjustment costs to include a nonconvex component. We also use this exercise to assess the accuracy of the log-linearized solution, since this case requires nonlinear solution techniques. We consider the implications of a frictional labor choice. We also explore a number of measurement concerns, including the potential for measurement error. Online Appendix I contains additional extensions and robustness exercises: alternative stochastic processes on productivity and distortions, different assumptions on the elasticity of substitution and variants on the set of target moments. Our main conclusions about the relative contribution of various factors to observed *arpk* dispersion are robust across these exercises.

A. Nonconvex Adjustment Costs

Our baseline specification with only convex adjustment costs allowed us to use perturbation techniques to solve and estimate the model, which yielded both analytical tractability for our identification arguments and computational efficiency. However, it raises two questions. One, how well does the log-linearized version approximate the true solution? And two, are the results robust to allowing for nonconvex adjustment costs? In this section, we address both of these concerns. We modify the adjustment cost function to include a nonconvex component:

$$\Phi(K_{it+1}, K_{it}) = \frac{\hat{\xi}}{2} \left(\frac{K_{it+1}}{K_{it}} - (1 - \delta) \right)^2 K_{it} + \hat{\xi}_f \mathbf{1}\{I_{it} \neq 0\} \pi(A_{it}, K_{it}),$$

where $I_{it} = K_{it+1} - (1 - \delta)K_{it}$ denotes period t investment and $\mathbf{1}\{\cdot\}$ the indicator function. The adjustment cost is now composed of two components: the first is a quadratic term, the same as before. The second is a fixed cost, which is parameterized by $\hat{\xi}_f$, the fraction of profits that must be incurred if the firm undertakes any nonzero investment.⁴⁶

⁴⁶The scaling with profits is a common formulation in the literature, see, e.g., Asker, Collard-Wexler, and De Loecker (2014), and ensures that the fixed cost does not become negligible for large firms.

TABLE 6—NONCONVEX ADJUSTMENT COSTS

Parameters	$\hat{\xi}$	(ξ)	$\hat{\xi}_f$	V	γ	σ_ε^2	σ_χ^2
China	0.075	(0.51)	0.002	0.09	−0.64	0.00	0.44
United States	0.250	(1.70)	0.002	0.03	−0.30	0.02	0.29
$\Delta\sigma_{arpk}^2/\sigma_{arpk}^2$							
China (%)	6.5		0.8	10.1	35.6	0.0	47.7
United States (%)	13.0		1.1	7.1	11.5	4.4	64.4

Notes: The second column (in parentheses) reports the value of the normalized adjustment cost parameter, ξ , for purposes of comparison to Table 3. The mapping between ξ and $\hat{\xi}$ is given in expression (A.1) in the online Appendix.

Due to the fixed cost, we can no longer use perturbation methods. We therefore solve the model by value function iteration and reestimate the parameters using simulated method of moments. Note that there is an additional parameter in this version relative to the baseline, $\hat{\xi}_f$. To pin this down, we add a new target moment: inaction, defined as the fraction of firms with (gross) investment rates of less than 5 percent in absolute value. This value is equal to 20 percent of firms in the Chinese data and 18 percent of firms in the United States.⁴⁷ Formally, we estimate the model by searching over the six parameters, $\hat{\xi}$, $\hat{\xi}_f$, V , γ , σ_ε^2 , and σ_χ^2 to find the combination that minimizes the equally-weighted sum of squared deviations of the model-implied values for the six target moments (the five moments from Table 2 and inaction) from their empirical counterparts.⁴⁸

The results are reported in Table 6. The estimated value for the fixed cost, $\hat{\xi}_f$, is modest in both countries, about 0.2 percent of annual profits. The other parameters and their contributions to σ_{arpk}^2 are quite close to the baseline estimates. These results demonstrate that (i) allowing for nonconvex adjustment costs does not alter our main conclusions regarding the sources of *arpk* dispersion and (ii) the perturbation approach produces reasonably accurate estimates.⁴⁹

Our estimates for the fixed adjustment cost are lower than many previous estimates in the literature.⁵⁰ The primary reason for this difference is the fact that we explicitly control for other factors and target a broader set of moments; indeed, our results here further underscore the importance of doing so. In online Appendix I.1, we explore this finding in greater depth by estimating a number of variants of our model with only adjustment costs, i.e., abstracting from other factors. This typically yields larger estimates of these costs, but at the expense of counterfactual implications

⁴⁷ Online Appendix I.1 shows the results are robust to using alternative moments to pin down the nonconvex component, for example, investment “spikes.”

⁴⁸ We provide further details of the estimation technique and results in online Appendix I.1; Table I.1 reports the fit of the model with respect to both targeted moments and nontargeted moments.

⁴⁹ Note that if $\hat{\xi} = \hat{\xi}_f = 0$, there is no approximation involved under the perturbation approach, i.e., the model is exactly log-linear. This property, along with the modest estimates for adjustment costs, is the main reason why the approximation works reasonably well in this region of the parameter space.

⁵⁰ For example, Bloom (2009) estimates a fixed adjustment cost of 1 percent of annual sales for US Compustat firms. Asker, Collard-Wexler, and De Loecker (2014) and Cooper and Haltiwanger (2006) work with data on US manufacturing firms and estimate this parameter at 12.5 percent of annual output and 4 percent of the capital stock, respectively.

for other, nontargeted moments. For example, a strategy which fits the low serial correlation of investment (e.g., Cooper and Haltiwanger 2006) yields much larger fixed costs (and smaller convex ones), but implies counterfactually high levels of investment variability and, even more strikingly, inaction. Importantly, however, even under this approach, adjustment costs generate only modest dispersion in *arpk*. Conversely, a strategy which matches the low variability of investment growth (along the lines of Asker, Collard-Wexler, and De Loecker 2014) implies larger convex costs but significantly over-predicts the autocorrelation of investment (and similarly leaves much of the *arpk* dispersion unexplained). Finally, a strategy that jointly fits both the serial correlation and variability yields larger costs of both types, but misses widely on other moments: for example, the predicted degree of inaction is extremely high relative to the data.

These patterns lead our estimation to ascribe an important role to other distortionary factors, even after allowing for nonconvexities: these factors reduce investment volatility without increasing its serial correlation and so reconcile these two moments. Further, in conjunction with the estimated level of adjustment costs, the model still performs well on additional moments such as inaction and investment spikes (see online Appendix I.1 for details).

B. Frictional Labor

Our baseline analysis makes the rather stark assumption of no adjustment or information frictions in labor choice, making it a static decision with full information. Although not uncommon in the literature, this may not be a good description of labor markets. Here, we depart from this assumption and assume that labor is subject to the same forces as capital: adjustment and informational frictions and other factors. In online Appendix E.2, we show that, under these conditions, the firm's investment problem takes the same form as in expression (3), but with a modified curvature parameter of $\alpha = \alpha_1 + \alpha_2$ (and appropriately redefined G and A_{it}). With this redefinition, our identification strategy goes through unchanged. Table 7 reports results for Chinese firms under this specification. The top panel shows the target moments recomputed under this assumption. A comparison to Table 2 reveals that assuming frictional labor raises the correlation of investment with lagged shocks as well as the correlation of the *arpk* with productivity. The second panel reports the associated parameter estimates. They imply higher adjustment costs, greater uncertainty and more severe correlated distortions. As a result, a lower level of the permanent factor, σ_{χ}^2 , is needed to match σ_{arpk}^2 .

The bottom panel of Table 7 reports the contribution of each factor to total *arpk* dispersion and computes the implications for aggregate TFP. There is a noticeable increase in the impact of adjustment costs from the baseline case: now, they account for almost 13 percent of *arpk* dispersion in China (compared to 1 percent above). There is also a slight increase in the impact of uncertainty (from 10 to 11 percent). Further, the effects on aggregate productivity are much larger than in the baseline scenario: here, these forces distort both inputs into production. Adjustment costs and imperfect information now lead to TFP losses of about 36 and 32 percent, respectively. Thus, this version of our model illustrates the potential for large aggregate consequences of adjustment/information frictions. However, despite the increased impact of these

TABLE 7—FRICTIONAL LABOR: CHINA

Moments	ρ	σ_μ^2	$\rho_{i,a_{-1}}$	$\rho_{i,k_{-1}}$	$\rho_{arpk,a}$	σ_i^2	σ_{arpk}^2
	0.92	0.16	0.33	-0.36	0.81	0.14	0.94
Parameters			ξ	V	γ	σ_ε^2	σ_χ^2
			0.78	0.11	-0.68	0.04	0.30
Aggregate effects							
$\Delta\sigma_{arpk}^2$			0.12	0.11	0.48	0.04	0.30
$\Delta\sigma_{arpk}^2/\sigma_{arpk}^2$			12.8%	11.3%	51.2%	4.0%	32.2%
Δa			0.36	0.32	1.44	0.11	0.90

forces (in both relative and absolute terms), the results also confirm a key finding from before, namely, the important role of other correlated and permanent factors. Indeed, these factors compose about 80 percent of the measured *arpk* dispersion, leading to TFP gaps relative to the first-best of about 144 and 90 percent, respectively.

C. Measurement Error

Measurement error is an important and challenging concern for the misallocation literature more broadly. In an important recent contribution, Bils, Klenow, and Ruane (2018) propose a method to estimate the role of additive measurement error. Here, we apply their methodology to our data. It essentially involves estimating the following regression:

$$\Delta va_{it} = \Phi arpk_{it} + \Psi \Delta k_{it} - \Psi(1 - \lambda) arpk_{it} \cdot \Delta k_{it} + D_{jt} + \epsilon_{it},$$

where Δva_{it} and Δk_{it} denote changes in (log) value-added and capital respectively, D_{jt} is a full set of industry-year fixed effects and $arpk_{it}$ is (the log of the) average revenue product of capital. The key object is the coefficient on the interaction term. Bils, Klenow, and Ruane (2018) show that, under certain assumptions, λ is the ratio of the true dispersion in the *arpk* to its measured counterpart (and inversely, $1 - \lambda$ is the contribution of measurement error to the observed σ_{arpk}^2). Intuitively, to the extent measured *arpk* deviations are due to additive measurement error, value-added of firms with high observed *arpk* will display a lower elasticity with respect to capital.

Estimating this regression in our data yields estimates for λ of 0.92 in China and 0.88 in the United States. These values suggest that, in both countries, only about 10 percent of the observed σ_{arpk}^2 can be accounted for by additive measurement error. Of course, it must be pointed out that this method is silent about other forms of measurement error (e.g., multiplicative).⁵¹

⁵¹ There are a few approaches in the literature to deal with multiplicative measurement error, e.g., Collard-Wexler and De Loecker (2016) and Song and Wu (2015) make some progress on this dimension after imposing additional structure.

D. Additional Measurement Concerns

In this subsection, we address two other measurement-related issues. The first stems from our use of book values for capital. Although this is a common approach in the misallocation literature, e.g., Hsieh and Klenow (2009) and Gopinath et al. (2017), other papers use the perpetual inventory method along with data on investment good price deflators to construct an alternative measure for capital. To address this concern, we compute firm-level capital stocks for US firms, where data on the relevant price indices are readily available, using the approach outlined in Eberly, Rebelo, and Vincent (2012). The results from reestimating the model using these measures, presented in online Appendix I.5, are broadly in line with our baseline findings. They point to a somewhat larger role for adjustment costs (the autocorrelation of investment growth is higher under this method and the variance lower, leading to a higher estimate of ξ), which account for about 27 percent of total σ_{arpk}^2 (compared to 11 percent under our baseline approach). The contribution of uncertainty is essentially unchanged at about 6 percent. Importantly, other firm-specific factors continue to play a key role in generating the observed *arpk* dispersion.

The second concern relates to sectoral heterogeneity in the structural parameters. We have estimated our model separately for US firms for the nine major sectors of the industrial classification (e.g., manufacturing, construction, services, etc.). Specifically, we allowed for sector-specific parameters in production (we infer sector-specific α 's using sectoral labor shares obtained from the Bureau of Economic Analysis), adjustment frictions, uncertainty, as well as other factors. The details of this procedure are outlined in online Appendix I.6 and the results are presented in online Appendix Table I.9. Although there is some variation across sectors, the overall patterns in the role of various factors (bottom panel of that table) are similar to those from our baseline analysis. The contribution of adjustment costs to observed *arpk* dispersion is generally modest: the highest contribution is about 20 percent of σ_{arpk}^2 in Manufacturing and the lowest is 2 percent in Finance, Insurance, and Real Estate. Uncertainty accounts for 5–10 percent across sectors, leaving the bulk of observed *arpk* dispersion within each sector to be accounted for by other factors.

VI. Conclusion

In this paper, we have laid out a model of investment featuring multiple factors that interfere with the equalization of static capital products, along with an empirical strategy to disentangle them using widely available firm-level production data. Figure 3 summarizes our results on the sources of *arpk* dispersion in China (panel A) and the United States (panel B). They show that much of the *arpk* dispersion stems not from adjustment and informational frictions, but from other firm-specific factors, either systematically correlated with firm productivity/size or almost permanent. Moreover, unobserved heterogeneity in demand and production technologies can potentially account for a significant portion of observed *arpk* dispersion in the United States, but not in China, where size-dependent policies and/or financial imperfections may be more fruitful avenues to pursue. Crucially, analyzing these forces in isolation would have led to very different conclusions, highlighting the value of using a unified framework and empirical approach.




FIGURE 3. THE SOURCES OF MISALLOCATION

Note: The numbers for the contribution of technological dispersion denote the upper bound as calculated in footnote 43.

There are several promising directions for future work. A key message of our analysis is that although various frictions/distortions/policies may all generate dispersion in input products, they often have different effects on various moments of the data, helping to tease out their individual contributions. This insight, along with our quantitative findings, should be useful both in identifying candidate factors driving that dispersion and in guiding empirical strategies to measure their impact. For example, policies and/or frictions that introduce persistent wedges into firms' investment decisions would be very promising; ones that are transitory, particularly those unrelated to firm size/productivity, less so. On the empirical side, our formulation and findings on these factors point to a strategy for investigating specific forces even while controlling for others, thereby reaching more accurate estimates of their contributions, in a tractable, albeit reduced-form, way. For example, one recent paper that builds on our results, both methodological and substantive, is David, Schmid, and Zeke (2018). First, they propose a theory of firm-level risk premia that delivers a firm-specific fixed wedge in the capital choice. Second, they verify that their identification strategy is robust to the presence of other distortions of the same form as we lay out here.

Our analysis focused primarily on capital allocation, but a natural extension is to use similar methods to study the allocation of labor. Such an analysis holds much promise, both for understanding the role of various forces (e.g., adjustment, informational, or other) in explaining observed dispersion in labor products, and further, in narrowing the list of candidate factors driving input allocations more broadly. As we showed in Section IVA, combining data on multiple inputs can help shed light on the nature of distortionary factors: for example, the correlation between capital and labor products can be very useful in disciplining the potential for "scale" factors that distort all input products in the same direction versus "mix" factors that distort the capital-labor ratio.

Our findings have implications beyond static *arpk* dispersion. Midrigan and Xu (2014) show that the same factors behind static variation in input products can have larger effects on aggregate outcomes by influencing entry and exit decisions. Similarly, a number of recent papers examine the impact of distortions on the life cycle of the firm and the distribution of productivity itself, e.g., Hsieh and Klenow (2014); Bento and Restuccia (2017); and Da-Rocha, Tavares, and Restuccia (2017). An important insight from these papers is that the exact nature of the underlying distortions (e.g., their correlation with firm productivity or demand) is key to understanding their dynamic implications. An ambitious next step would be to use an empirical strategy like the one in this paper to analyze richer environments featuring some of these elements.

REFERENCES

- Asker, John, Allan Collard-Wexler, and Jan De Loecker. 2014. "Dynamic Inputs and Resource (Mis)Allocation." *Journal of Political Economy* 122 (5): 1013–63.
- Bachmann, Ruediger, and Steffen Elstner. 2015. "Firm Optimism and Pessimism." *European Economic Review* 79: 297–325.
- Bai, Chong-En, Chang-Tai Hsieh, and Yingyi Qian. 2006. "The Return to Capital in China." *Brookings Papers on Economic Activity* 2: 61–101.
- Bartelsman, Eric, John Haltiwanger, and Stefano Scarpetta. 2013. "Cross-Country Differences in Productivity: The Role of Allocation and Selection." *American Economic Review* 103 (1): 305–34.
- Bento, Pedro, and Diego Restuccia. 2017. "Misallocation, Establishment Size, and Productivity." *American Economic Journal: Macroeconomics* 9 (3): 267–303.
- Bils, Mark, Peter J. Klenow, and Cian Ruane. 2018. "Misallocation or Mismeasurement?" <http://klenow.com/misallocation-mismeasurement-paper.pdf>.
- Bloom, Nicholas. 2009. "The Impact of Uncertainty Shocks." *Econometrica* 77 (3): 623–85.
- Brandt, Loren, Johannes Van Biesebroeck, and Yifan Zhang. 2014. "Challenges of Working with the Chinese NBS Firm-Level Data." *China Economic Review* 30: 339–52.
- Buera, Francisco J., and Roberto N. Fattal-Jaef. 2018. "The Dynamics of Development: Innovation and Reallocation?" https://www.dropbox.com/s/5jyc5u7xlyi5u4u/draft_world_bank_june_2018.pdf.
- Buera, Francisco J., Joseph P. Kaboski, and Yongseok Shin. 2011. "Finance and Development: A Tale of Two Sectors." *American Economic Review* 101 (5): 1964–2002.
- Buera, Francisco J., Benjamin Moll, and Yongseok Shin. 2013. "Well-Intended Policies." *Review of Economic Dynamics* 16 (1): 216–30.
- Collard-Wexler, Allan, and Jan De Loecker. 2016. "Production Function Estimation with Measurement Error in Inputs." NBER Working Paper 22437.
- Cooper, Russell, Guan Gong, and Ping Yan. 2015. "Dynamic Labor Demand in China: Public and Private Objectives." *RAND Journal of Economics* 46 (3): 577–610.
- Cooper, Russell W., and John C. Haltiwanger. 2006. "On the Nature of Capital Adjustment Costs." *Review of Economic Studies* 73 (3): 611–33.
- Da-Rocha, José-María, Marina Mendes Tavares, and Diego Restuccia. 2017. "Policy Distortions and Aggregate Productivity with Endogenous Establishment-Level Productivity." NBER Working Paper 23339.
- David, Joel M., Hugo A. Hopenhayn, and Venky Venkateswaran. 2016. "Information, Misallocation, and Aggregate Productivity." *Quarterly Journal of Economics* 131 (2): 943–1005.
- David, Joel M., Lukas Schmid, and David Zeke. 2018. "Risk-Adjusted Capital Allocation and Misallocation." https://sites.google.com/site/joelmichaeldavid/misalloc_risk.pdf.
- David, Joel M., and Venky Venkateswaran. 2019. "The Sources of Capital Misallocation: Dataset." *American Economic Review*. <https://doi.org/10.1257/aer.20180336>.
- De Loecker, Jan, and Jan Eeckhout. 2017. "The Rise of Market Power and the Macroeconomic Implications." NBER Working Paper 23687.
- De Loecker, Jan, and Frederic Warzynski. 2012. "Markups and Firm-Level Export Status." *American Economic Review* 102 (6): 2437–71.
- Eberly, Janice, Sérgio Rebelo, and Nicolas Vincent. 2012. "What Explains the Lagged-Investment Effect?" *Journal of Monetary Economics* 59 (4): 370–80.

- Gopinath, Gita, Şebnem Kalemli-Özcan, Loukas Karabarbounis, and Carolina Villegas-Sanchez.** 2017. "Capital Allocation and Productivity in South Europe." *Quarterly Journal of Economics* 132 (4): 1915–67.
- Guner, Nezih, Gustavo Ventura, and Yi Xu.** 2008. "Macroeconomic Implications of Size-Dependent Policies." *Review of Economic Dynamics* 11 (4): 721–44.
- Hopenhayn, Hugo A.** 2014. "Firms, Misallocation, and Aggregate Productivity: A Review." *Annual Review of Economics* 6: 735–70.
- Hsieh, Chang-Tai, and Peter J. Klenow.** 2009. "Misallocation and Manufacturing TFP in China and India." *Quarterly Journal of Economics* 124 (4): 1403–48.
- Hsieh, Chang-Tai, and Peter J. Klenow.** 2014. "The Life Cycle of Plants in India and Mexico." *Quarterly Journal of Economics* 129 (3): 1035–84.
- İmrohoroglu, Ayşe, and Şelale Tüzel.** 2014. "Firm-Level Productivity, Risk, and Return." *Management Science* 60 (8): 2073–90.
- Jurado, Kyle, Sydney C. Ludvigson, and Serena Ng.** 2015. "Measuring Uncertainty." *American Economic Review* 105 (3): 1177–216.
- Kehrig, Matthias, and Nicolas Vincent.** 2017. "Do Firms Mitigate or Magnify Capital Misallocation? Evidence from Planet-Level Data." https://papers.ssrn.com/sol3/papers.cfm?abstract_id=2731594.
- Khan, Aubhik, and Julia K. Thomas.** 2008. "Idiosyncratic Shocks and the Role of Nonconvexities in Plant and Aggregate Investment Dynamics." *Econometrica* 76 (2): 395–436.
- Klenow, Peter J., and Jonathan L. Willis.** 2007. "Sticky Information and Sticky Prices." *Journal of Monetary Economics* 54: S79–99.
- Midrigan, Virgiliu, and Daniel Yi Xu.** 2014. "Finance and Misallocation: Evidence from Plant-Level Data." *American Economic Review* 104 (2): 422–58.
- Moll, Benjamin.** 2014. "Productivity Losses from Financial Frictions: Can Self-Financing Undo Capital Misallocation?" *American Economic Review* 104 (10): 3186–221.
- Morck, Randall, Andrei Shleifer, and Robert W. Vishny.** 1990. "The Stock Market and Investment: Is the Market a Sideshow?" *Brookings Papers on Economic Activity* 2: 157–202.
- Peters, Michael.** 2016. "Heterogeneous Markups, Growth and Endogenous Misallocation." <https://mipeters.weebly.com/uploads/1/4/6/5/14651240/markupsmisallocation.pdf>.
- Restuccia, Diego, and Richard Rogerson.** 2008. "Policy Distortions and Aggregate Productivity with Heterogeneous Establishments." *Review of Economic Dynamics* 11 (4): 707–20.
- Restuccia, Diego, and Richard Rogerson.** 2017. "The Causes and Costs of Misallocation." *Journal of Economic Perspectives* 31 (3): 151–74.
- Song, Zheng, and Guiying Laura Wu.** 2015. "Identifying Capital Misallocation." https://michaelzsong.weebly.com/uploads/4/8/1/4/48141215/sw_web.pdf.